


BİLGİSAYAR MİMARİSİ

Bellek Yönetim Sistemi

Bölüm 7

- **Memory Hierarchy**
- **Associative Memory**
- **Cache Memory**
- **Virtual Memory**
- **Memory Management Hardware**

An Expanded View of the Memory System



The important characteristics of a device:

- **Access Time**

- **average memory access time (AMAT)** is a common metric to analyze memory system performance. AMAT uses hit time, miss penalty, and miss rate to measure memory performance.

- **Transfer Rate**

- The memory transfer rate is determined by three factors; the memory bus clock rate, the type of transfer process and the number of bits transferred.
- **Step 1**
- Determine the bus clock rate. For example, the memory might operate at 300 MHz.
- **Step 2**
- Multiply the bus **clock** rate by 1, 2 or 3 depending on how many streams of information are flowing at once. DDR2 RAM, for example, has a multiplication factor of 2, and DDR3 RAM has a factor of 3.
- **Step 3**
- Multiply the result from the previous step by 64, which is the number of bits transferred.
- Divide the result from the previous step by 8 to get the transfer rate in bytes instead of bits, as there are 8 bits in a byte. You now know the memory transfer rate in both bits and bytes
- **Capacity**
- The total memory (RAM) that can be added to a computer depends on the address registers built into the CPU. For example, most 32-bit CPUs can address only up to 4 gigabytes (GB) of memory.

Cost

Levels of the Memory Hierarchy

Capacity
Access Time
Cost

CPU Registers
100s Bytes
<10s ns

Cache
K Bytes
10-100 ns
\$.01-.001/bit

Main Memory
M Bytes
100ns-1us
\$.01-.001

Disk
G Bytes
ms₃ - 10⁻⁴
10⁻³ - 10 cents

Tape
infinite
sec-min
10⁻⁶

Registers

Instr. Operands

Cache

Blocks

Memory

Pages

Disk

Files

Tape

Staging
Xfer Unit

prog./compiler
1-8 bytes

cache cntl
8-128 bytes

OS
512-4K bytes

user/operator
Mbytes

Upper Level

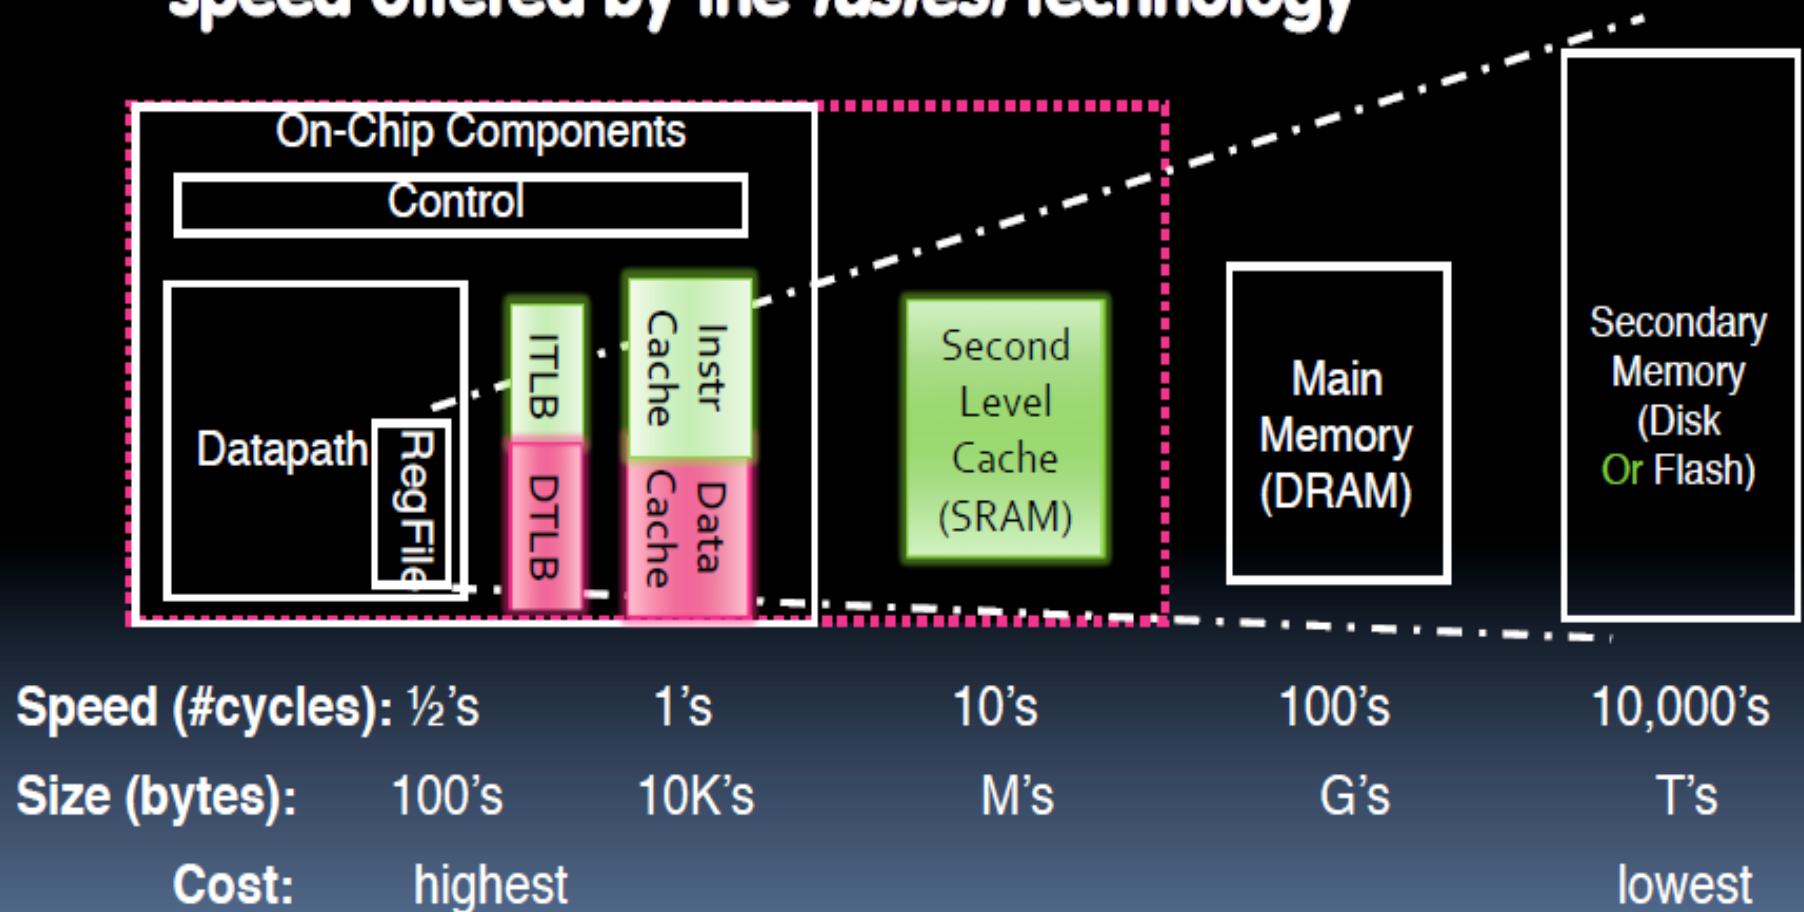
faster

Lower Level

Larger

Typical Memory Hierarchy

- The Trick:** present processor with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology



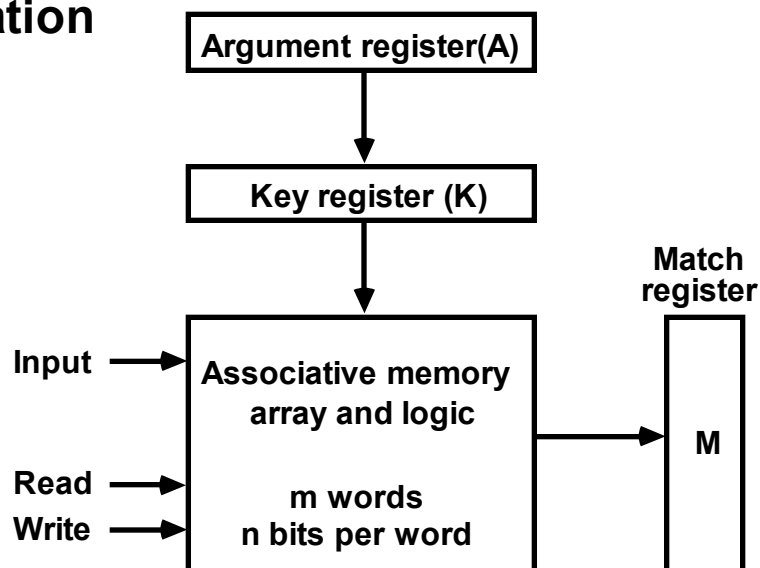
Memory Hierarchy Analogy: Library

- You're writing a term paper (Processor) at a **table** in **Doe**
- **Doe Library** is equivalent to **disk**
 - essentially limitless capacity, very slow to retrieve a book
- **Table** is **main memory**
 - smaller capacity: means you must return book when table fills up
 - easier and faster to find a book there once you've already retrieved it
- Open books on table are **cache**
 - smaller capacity: can have very few open books fit on table; again, when table fills up, you must close a book
 - much, much faster to retrieve data
- Illusion created: whole library open on the tabletop
 - Keep as many recently used books open on table as possible since likely to use again
 - Also keep as many books on table as possible, since faster than going to library

ASSOCIATIVE MEMORY

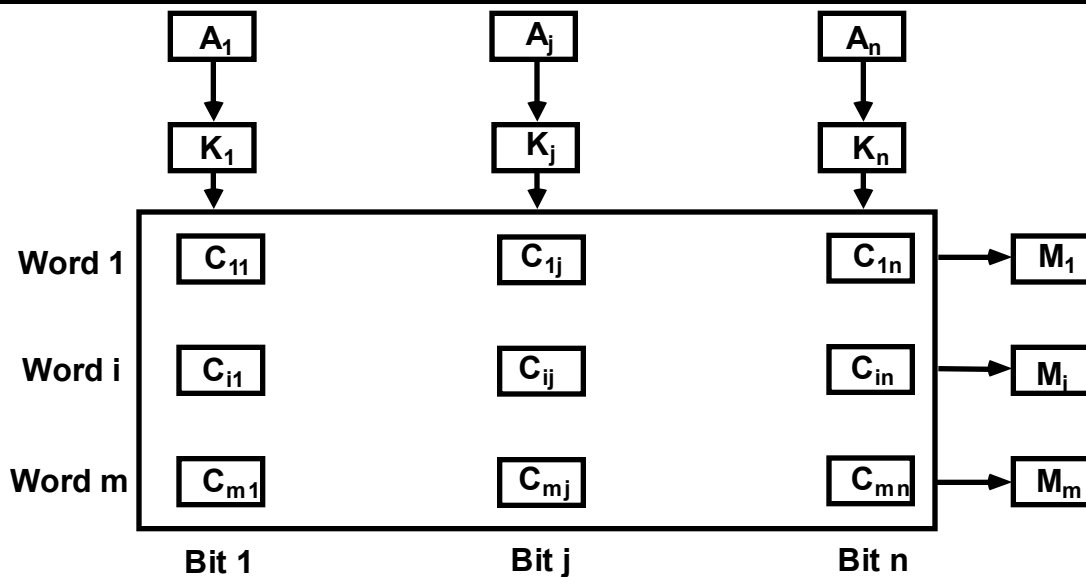
- Accessed by the content of the data rather than by an address
- Also called Content Addressable Memory (CAM)

Hardware Organization

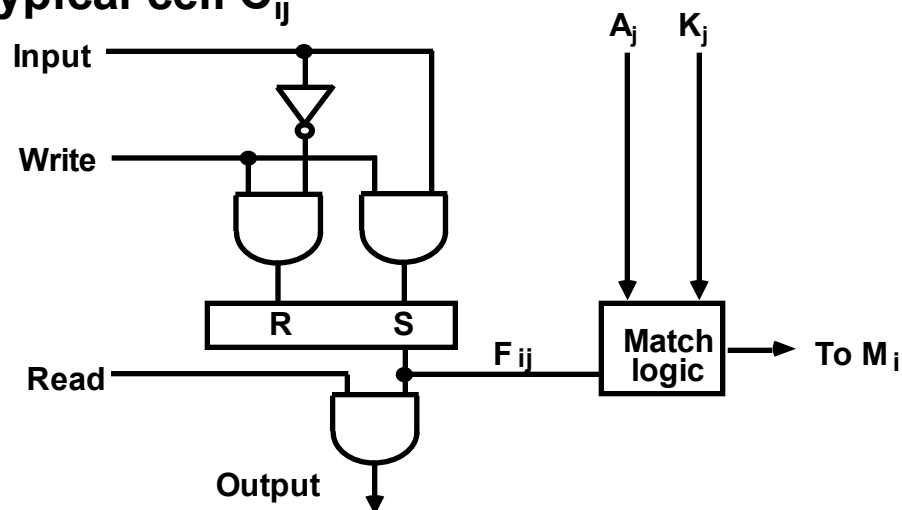


- Compare each word in CAM in parallel with the content of A(Argument Register)
- If CAM Word[i] = A, M(i) = 1
- Read sequentially accessing CAM for CAM Word(i) for M(i) = 1
- K(Key Register) provides a mask for choosing a particular field or key in the argument in A (only those bits in the argument that have 1's in their corresponding position of K are compared)


ORGANIZATION OF CAM



Internal organization of a typical cell C_{ij}



MATCH LOGIC



CACHE MEMORY


Locality of Reference

- The references to memory at any given time interval tend to be confined within a localized areas
- This area contains a set of information and the membership changes gradually as time goes by
- *Temporal Locality*
The information which will be used in near future is likely to be in use already
(e.g. Reuse of information in loops)
- *Spatial Locality*
If a word is accessed,
adjacent(near) words are likely accessed soon

(e.g. Related data items (arrays) are usually stored together;
instructions are executed sequentially)


Terminology

- **Hit:** data appears in some block in the upper level (example: Block X)
 - Hit Rate: the fraction of memory access found in the upper level
 - Hit Time: Time to access the upper level which consists of
RAM access time + Time to determine hit/miss
- **Miss:** data needs to be retrieve from a block in the lower level (Block Y)
 - Miss Rate = $1 - (\text{Hit Rate})$
 - Miss Penalty = Time to replace a block in the upper level +
Time to deliver the block the processor
- **Hit Time \ll Miss Penalty**



Principles of Operation

- **At any given time, data is copied between only 2 adjacent levels:**
 - **Upper Level (Cache) :** the one closer to the processor
 - » Smaller, faster, and uses more expensive technology
 - **Lower Level (Memory):** the one further away from the processor
 - » Bigger, slower, and uses less expensive technology
- **Block:**
 - The minimum unit of information that can either be present or not present in the two level hierarchy



Typical Values

Block (line) size 4 - 128 bytes

Hit time 1 - 4 cycles


Miss penalty 8 - 32 cycles (and increasing)
 (access time) (6-10 cycles)
 (transfer time) (2 - 22 cycles)

Miss rate 1% - 20%

Cache Size 1 KB - 256 KB

How Does Cache Work?

- **Temporal Locality (Locality in Time):** If an item is referenced, it will tend to be referenced again soon.
 - Keep more recently accessed data items closer to the processor
- **Spatial Locality (Locality in Space):** If an item is referenced, items whose addresses are close by tend to be referenced soon.
 - Move blocks consists of contiguous words to the cache



Memory Access with Cache

- Load word instruction: lw \$t0, 0(\$t1)
- \$t1 contains 1022_{ten} , Memory[1022] = 99
- With cache (similar to a hash)
 1. Processor issues address 1022_{ten} to Cache
 2. Cache checks to see if has copy of data at address 1022_{ten}
 - 2a. If finds a match (Hit): cache reads 99, sends to processor
 - 2b. No match (Miss): cache sends address 1022 to Memory
 - I. Memory reads 99 at address 1022_{ten}
 - II. Memory sends 99 to Cache
 - III. Cache replaces word with new 99
 - IV. Cache sends 99 to processor
 3. Processor loads 99 into register \$t1

PERFORMANCE OF CACHE

Performance of Cache Memory System

Hit Ratio - % of memory accesses satisfied by Cache memory system

Te: Effective memory access time in Cache memory system

Tc: Cache access time

Tm: Main memory access time

$$T_e = T_c * h + (1 - h) T_m$$

Example: $T_c = 0.4 \mu s$, $T_m = 1.2 \mu s$, $h = 0.85\%$

$$T_e = 0.4 * 0.85 + (1 - 0.85) * 1.2 = 0.52 \mu s$$

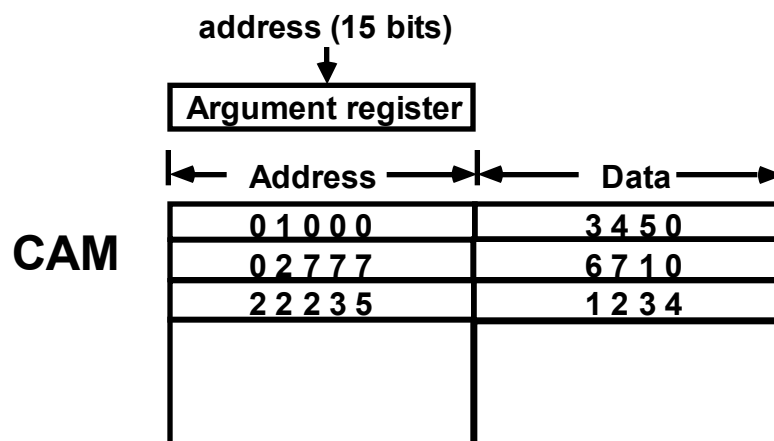
Mapping Function

**Specification of correspondence
between main memory blocks and cache blocks**

- **Associative mapping**
- **Direct mapping**
- **Set-associative mapping**

ASSOCIATIVE MAPPING -


- Any block location in Cache can store any block in memory
-> Most flexible
- Mapping Table is implemented in an associative memory
-> Fast, very Expensive
- Mapping Table
Stores both address and the content of the memory word



DIRECT MAPPING -

- Each memory block has only one place to load in Cache
- Mapping Table is made of RAM instead of CAM
- n-bit memory address consists of 2 parts; k bits of Index field and n-k bits of Tag field
- n-bit addresses are used to access main memory and k-bit Index is used to access the Cache

Addressing Relationships



Direct Mapping Cache Organization

Memory address	Memory data
00000	1 2 2 0
00777	2 3 4 0
01000	3 4 5 0
01777	4 5 6 0
02000	5 6 7 0
02777	6 7 1 0

Cache memory		
Index address	Tag	Data
000	0 0	1 2 2 0
777	0 2	6 7 1 0

The Simplest Cache: Direct Mapped Cache

Memory Address Memory

0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
A	
B	
C	
D	
E	
F	

4 Byte Direct Mapped Cache

Cache Index

0	
1	
2	
3	

- Location 0 can be occupied by data from:
 - Memory location 0, 4, 8, ... etc.
 - In general: any memory location whose 2 LSBs of the address are 0s
 - $\text{Address} \langle 1:0 \rangle \Rightarrow \text{cache index}$
- Which one should we place in the cache?
- How can we tell which one is in the cache?


DIRECT MAPPING

Operation

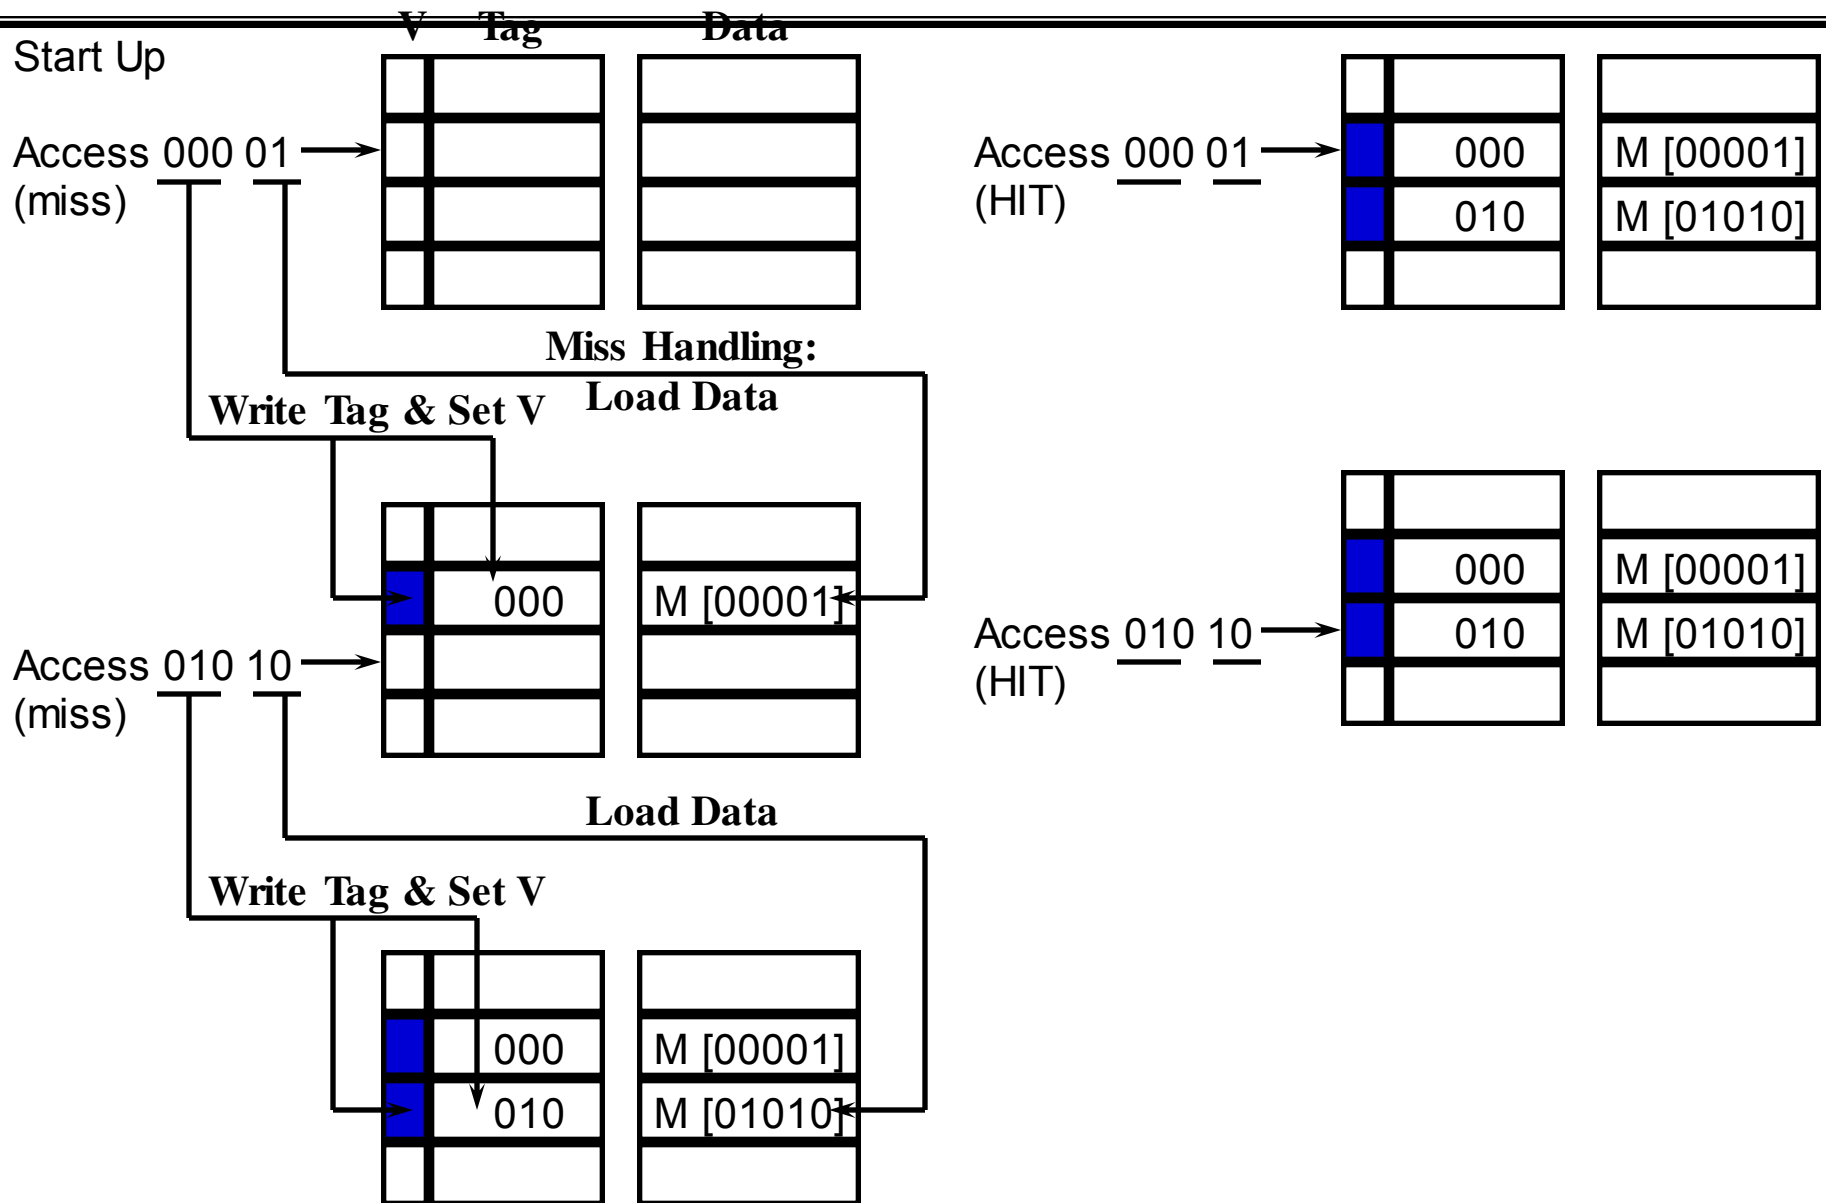
- CPU generates a memory request with (TAG;INDEX)
- Access Cache using INDEX ; (tag; data)
 - Compare TAG and tag
- If matches -> Hit
 - Provide Cache[INDEX](data) to CPU
- If not match -> Miss
 - $M[\text{tag}; \text{INDEX}] \leftarrow \text{Cache}[\text{INDEX}](\text{data})$
 - $\text{Cache}[\text{INDEX}] \leftarrow (\text{TAG}; M[\text{TAG}; \text{INDEX}])$
 - $\text{CPU} \leftarrow \text{Cache}[\text{INDEX}](\text{data})$

Direct Mapping with block size of 8 words

	Index	tag	data
Block 0	000	0 1	3 4 5 0
	007	0 1	6 5 7 8
Block 1	010		
	017		
...			
Block 63	770	0 2	
	777	0 2	6 7 1 0




Cache Access Example



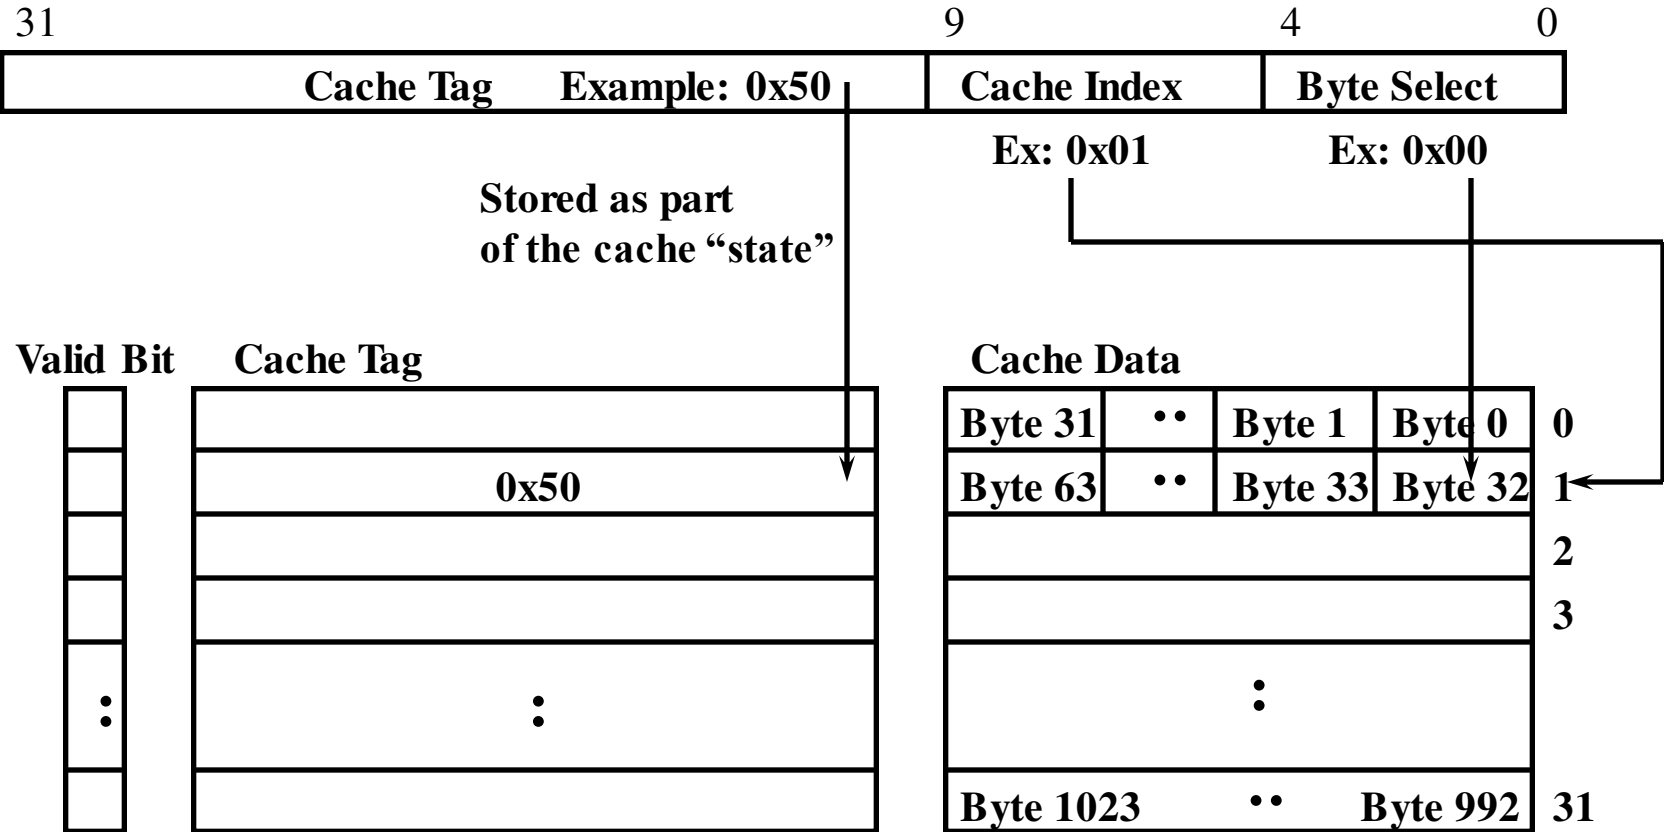
Definition of a Cache Block

- **Cache Block:** the cache data that has in its own cache tag
 - **4-byte Direct Mapped Cache:** Block Size = 1 Byte
 - **Take advantage of Temporal Locality:** If a byte is referenced, it will tend to be referenced soon.
 - **Did not take advantage of Spatial Locality:** If a byte is referenced, its adjacent bytes will be referenced soon.
- **In order to take advantage of Spatial Locality:** increase the block size



Ex: 1 KB Direct Mapped Cache with 32 B Blocks

- For a 2^N byte cache:
 - The uppermost $(32 - N)$ bits are always the Cache Tag
 - The lowest M bits are the Byte Select (Block Size = 2^M)



Accessing data in a direct mapped cache

- Ex.: 16KB of data, direct-mapped, 4 word blocks

- Can you work out height, width, area?

- Read 4 addresses

1. 0x00000014
2. 0x0000001C
3. 0x00000034
4. 0x00008014

- Memory vals here:

Memory
Address (hex) Value of Word

...
00000010
00000014
00000018
0000001C
...

...
a
b
c
d

...
00000030
00000034
00000038
0000003C
...

...
e
f
g
h

...
00008010
00008014
00008018
0000801C

...
i
j
k
l

Accessing data in a direct mapped cache

- **4 Addresses:**

- 0x00000014, 0x0000001C,
0x00000034, 0x00008014

- **4 Addresses divided (for convenience) into Tag, Index, Byte Offset fields**

00000000000000000000 0000000001 0100

00000000000000000000 0000000001 1100

00000000000000000000 0000000011 0100

00000000000000000010 0000000001 0100

Tag

Index

Offset

16 KB Direct Mapped Cache, 16B blocks

- **Valid bit:** determines whether anything is stored in that row (when computer initially turned on, all entries invalid)

<u>Valid</u>						
Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3	
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					

1. Read 0x00000014

▪ 000000000000000000000000 00000000001 0100

Valid Tag field Index field Offset

Index	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					

...

...

1022	0					
1023	0					

So we read block 1 (0000000001)

▪ 000000000000000000000000 0000000001 0100

Tag field Index field Offset

Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index					
0	0				
1	0				
2	0				
3	0				
4	0				
5	0				
6	0				
7	0				

...

...

1022	0				
1023	0				

So load that data into cache, setting tag, valid

- 00000000000000000000 0000000001 0100

Valid		Tag field		Index field		Offset
Index		Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...		...				
1022	0					
1023	0					

So read block 3

▪ 000000000000000000000000 0000000011 0100
Tag field Index field Offset

Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index	Tag				
0	0				
1	1	d	c	b	a
2	0				
3	0				
4	0				
5	0				
6	0				
7	0				

...

...

1022	0				
1023	0				

Load that cache block, return word f

- 000000000000000000000000 00000000011 0100
Tag field Index field Offset

Valid	Index	Tag	0xc-f	0x8-b	<u>0x4-7</u>	0x0-3
0	0					
1	1	0	d	c	b	a
0	2					
1	3	0	h	g	f	e
0	4					
0	5					
0	6					
0	7					

...

...

1022	0					
1023	0					

4. Read 0x00008014 = 0...10 0..001 0100

▪ 0000000000000000000010 0000000001 0100

Valid Tag field Index field Offset

Index	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					

...

...

1022	0					
1023	0					

So read Cache Block 1, Data is Valid

▪ 0000000000000000000010 0000000001 0100
Tag field Index field Offset

Valid

Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0				
<u>1</u>	0	d	c	b	a
2	0				
3	1	h	g	f	e
4	0				
5	0				
6	0				
7	0				

...

...

1022	0				
1023	0				

Cache Block 1 Tag does not match (0 != 2)

- 0000000000000000000010 000000000001 0100
Tag field Index field Offset

Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index					
0	0				
1	0	d	c	b	a
2	0				
3	1	h	g	f	e
4	0				
5	0				
6	0				
7	0				

...

...

1022	0				
1023	0				

Miss, so replace block 1 with new data & tag

- 0000000000000000000010 000000000001 0100

Valid		Tag field		Index field		Offset
Index	Tag	0xc-f	0x8-b	0x4-7	0x0-3	
0	0					
1	1	2	l	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...		...				
1022	0					
1023	0					

And return word J

▪ 0000000000000000000010 0000000001 0100

Valid		Tag field		Index field		Offset
Index		Tag	0xc-f	0x8-b	<u>0x4-7</u>	0x0-3
0	0					
1	1	2	l	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					




...

...

1022	0					
1023	0					

Block Size Tradeoff

- In general, larger block size take advantage of spatial locality BUT:
 - Larger block size means larger miss penalty:
 - » Takes longer time to fill up the block
 - If block size is too big relative to cache size, miss rate will go up
- Average Access Time:
 - $\text{Hit Time} \times (1 - \text{Miss Rate}) + \text{Miss Penalty} \times \text{Miss Rate}$



SET ASSOCIATIVE MAPPING -

- Each memory block has a set of locations in the Cache to load

Set Associative Mapping Cache with set size of two


Index	Tag	Data	Tag	Data
000	0 1	3 4 5 0	0 2	5 6 7 0
777	0 2	6 7 1 0	0 0	2 3 4 0

Operation

- CPU generates a memory address(TAG; INDEX)
- Access Cache with INDEX, (Cache word = (tag 0, data 0); (tag 1, data 1))
- Compare TAG and tag 0 and then tag 1
- If tag i = TAG -> Hit, CPU <- data i
- If tag i ≠ TAG -> Miss,
 Replace either (tag 0, data 0) or (tag 1, data 1),
 Assume (tag 0, data 0) is selected for replacement,
 M[tag 0, INDEX] <- Cache[INDEX](data 0)
 Cache[INDEX](tag 0, data 0) <- (TAG, M[TAG,INDEX]),
 CPU <- Cache[INDEX](data 0)

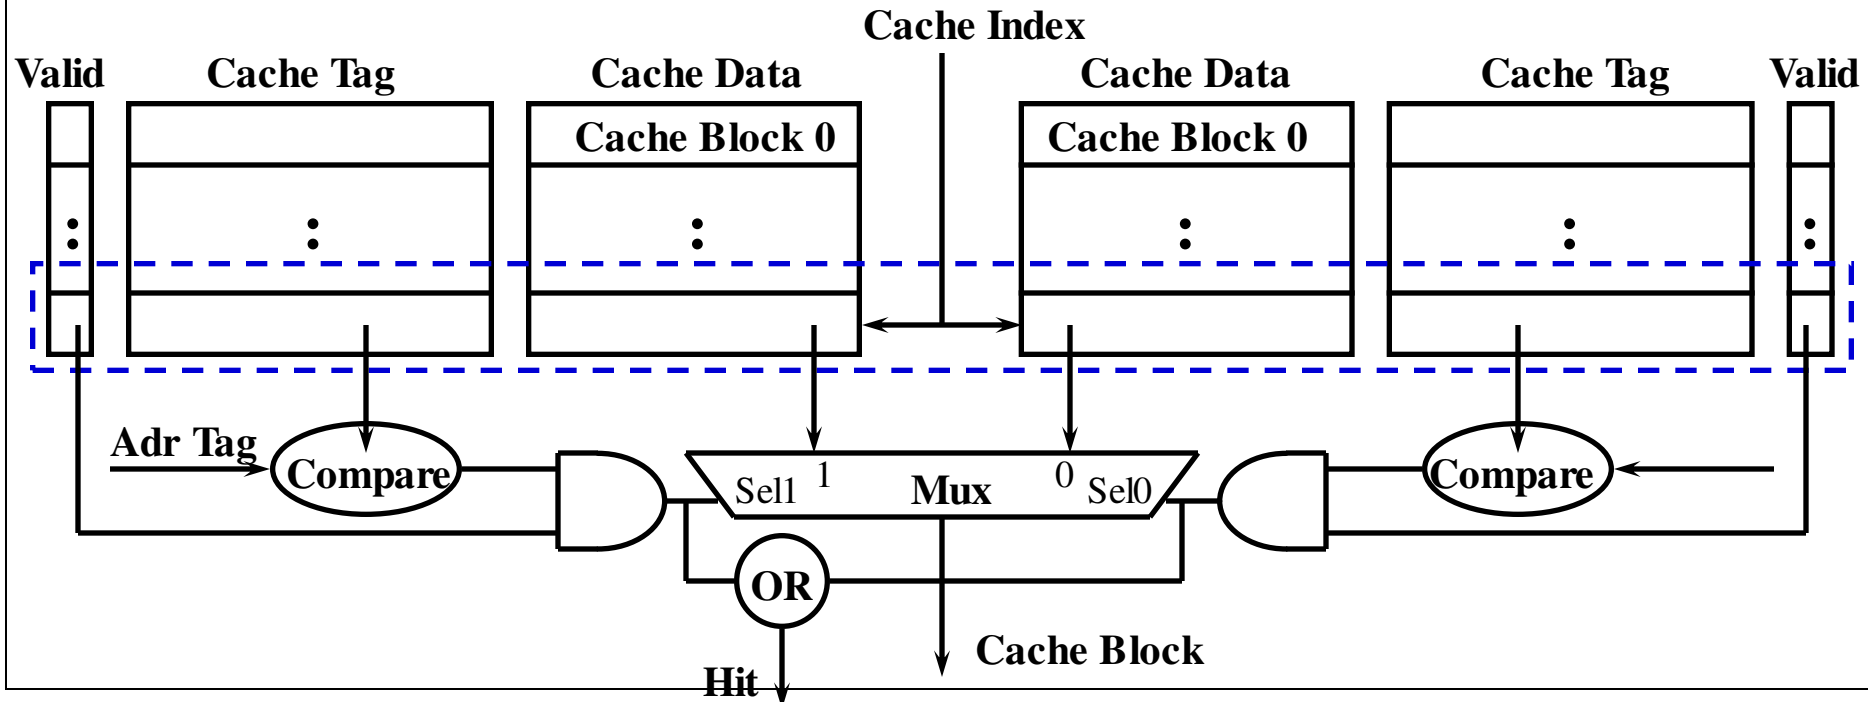
A Two-way Set Associative Cache

- **N-way set associative: N entries for each Cache Index**
 - N direct mapped caches operates in parallel
- **Example: Two-way set associative cache**
 - Cache Index selects a “set” from the cache
 - The two tags in the set are compared in parallel
 - Data is selected based on the tag result




Disadvantage of Set Associative Cache

- **N-way Set Associative Cache versus Direct Mapped Cache:**
 - N comparators vs. 1
 - Extra MUX delay for the data
 - Data comes AFTER Hit/Miss
- **In a direct mapped cache, Cache Block is available BEFORE Hit/Miss:**
 - Possible to assume a hit and continue.



And yet Another Extreme Example: Fully Associative

- ° Fully Associative Cache -- push the set associative idea to its limit!
 - Forget about the Cache Index
 - Compare the Cache Tags of all cache entries in parallel
 - Example: Block Size = 2 B blocks, we need N 27-bit comparators
- ° By definition: Conflict Miss = 0 for a fully associative cache



The Need to Make a Decision!

- Direct Mapped Cache:
 - Each memory location can only mapped to 1 cache location
 - No need to make any decision :-)
 - Current item replaced the previous item in that cache location
- N-way Set Associative Cache:
 - Each memory location have a choice of N cache locations
- Fully Associative Cache:
 - Each memory location can be placed in ANY cache location
- Cache miss in a N-way Set Associative or Fully Associative Cache:
 - Bring in new block from memory
 - Throw out a cache block to make room for the new block
 - Damn! We need to make a decision on which block to throw out!

Replacement Algorithms

- *When a block is fetched, which block in the target set should be replaced?*
- Optimal algorithm:
 - | replace the block that will not be used for the longest time (must know the future)
- Usage based algorithms:
 - Least recently used (LRU)
 - | replace the block that has been referenced least recently
 - | hard to implement
- Non-usage based algorithms:
 - First-in First-out (FIFO)
 - | treat the set as a circular queue, replace head of queue.
 - | easy to implement
 - Random (RAND)
 - | replace a random block in the set
 - | even easier to implement

LRU(Least Recently Used)

Cache word = (tag 0, data 0, U_0);(tag 1, data 1, U_1), $U_i = 0$ or 1(binary)

Implementation of LRU in the Set Associative Mapping with set size = 2

Modifications

Initially all $U_0 = U_1 = 1$

When Hit to (tag 0, data 0, U_0), $U_1 \leftarrow 1$ (least recently used)
(When Hit to (tag 1, data 1, U_1), $U_0 \leftarrow 1$ (least recently used))

When Miss, find the least recently used one($U_i=1$)

If $U_0 = 1$, and $U_1 = 0$, then replace (tag 0, data 0)

$M[\text{tag } 0, \text{INDEX}] \leftarrow \text{Cache}[\text{INDEX}](\text{data } 0)$

$\text{Cache}[\text{INDEX}](\text{tag } 0, \text{data } 0, U_0) \leftarrow (\text{TAG}, M[\text{TAG}, \text{INDEX}], 0); U_1 \leftarrow 1$

If $U_0 = 0$, and $U_1 = 1$, then replace (tag 1, data 1)

Similar to above; $U_0 \leftarrow 1$

If $U_0 = U_1 = 0$, this condition does not exist

If $U_0 = U_1 = 1$, Both of them are candidates,

Take arbitrary selection

CACHE WRITE

Write Through

When writing into memory

If Hit, both Cache and memory is written in parallel

If Miss, Memory is written

For a read miss, missing block may be overloaded onto a cache block

Memory is always updated

-> Important when CPU and DMA I/O are both executing

Slow, due to the memory access time

Write-Back (Copy-Back)

When writing into memory

If Hit, only Cache is written


If Miss, missing block is brought to Cache and write into Cache

For a read miss, candidate block must be written back to the memory

Memory is not up-to-date, i.e., the same item in Cache and memory may have different value


Write Through

- Store by processor updates cache *and* memory
- Memory always consistent with cache
- ~2X more loads than stores
- WT always combined with write buffers so that don't wait for lower level memory

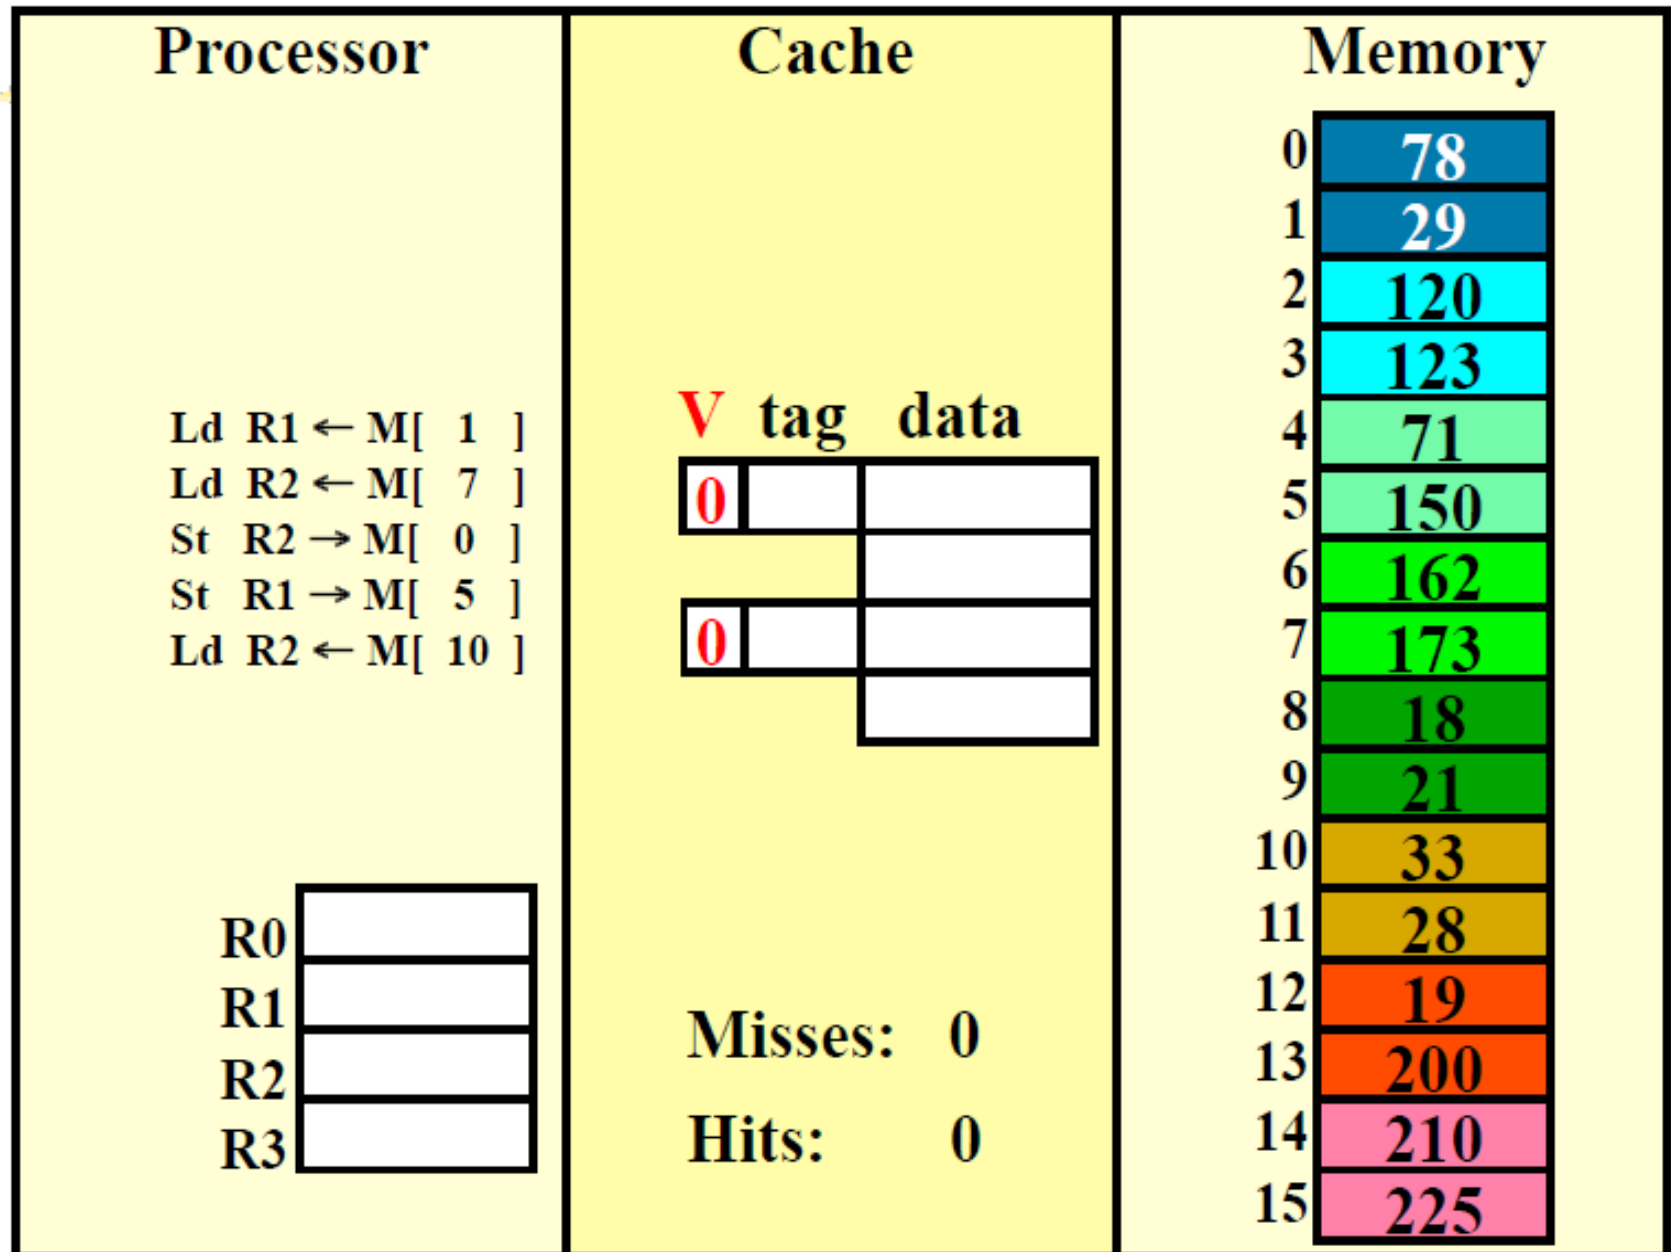


Write Back

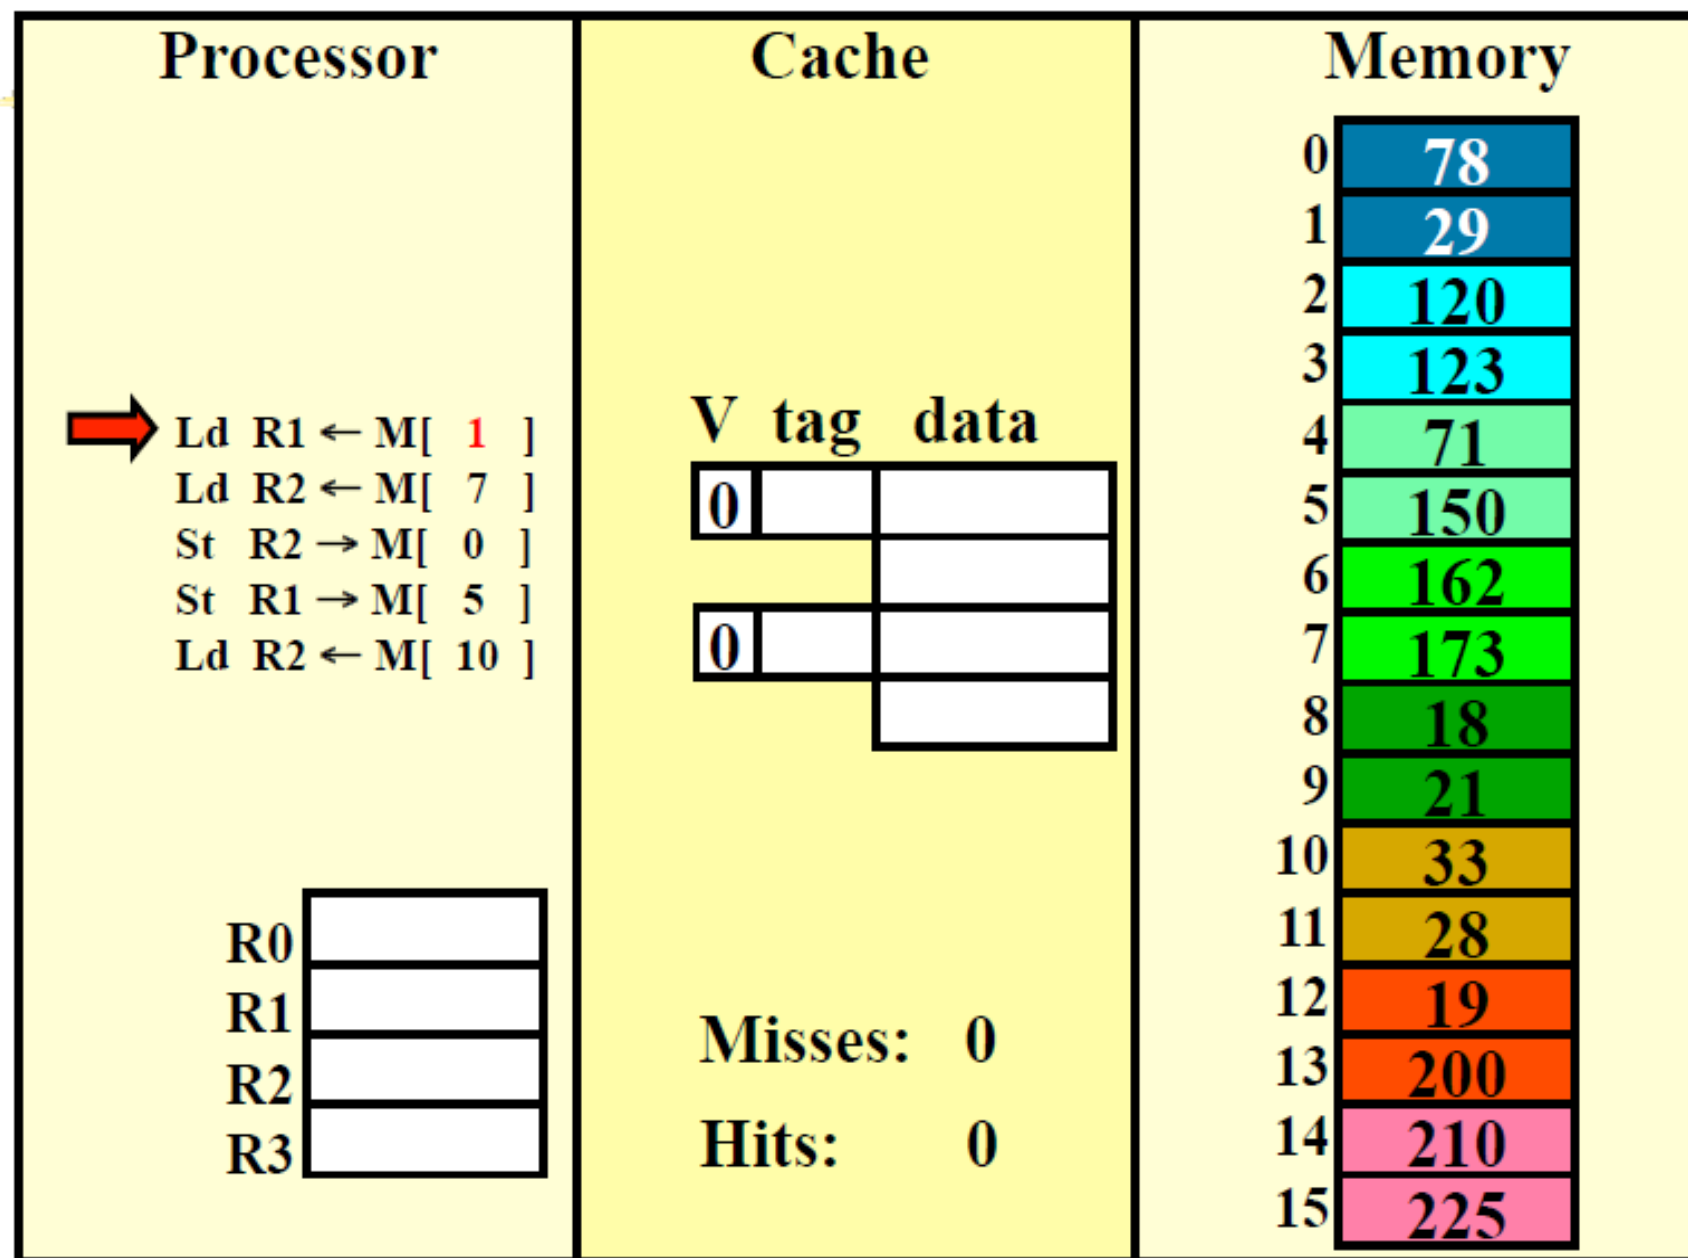
- Store by processor only updates cache line
- Modified line written to memory only when it is evicted
 - Requires “dirty bit” for each line
 - Set when line in cache is modified
 - Indicates that line in memory is stale
- Memory not always consistent with cache
- No writes of repeated writes



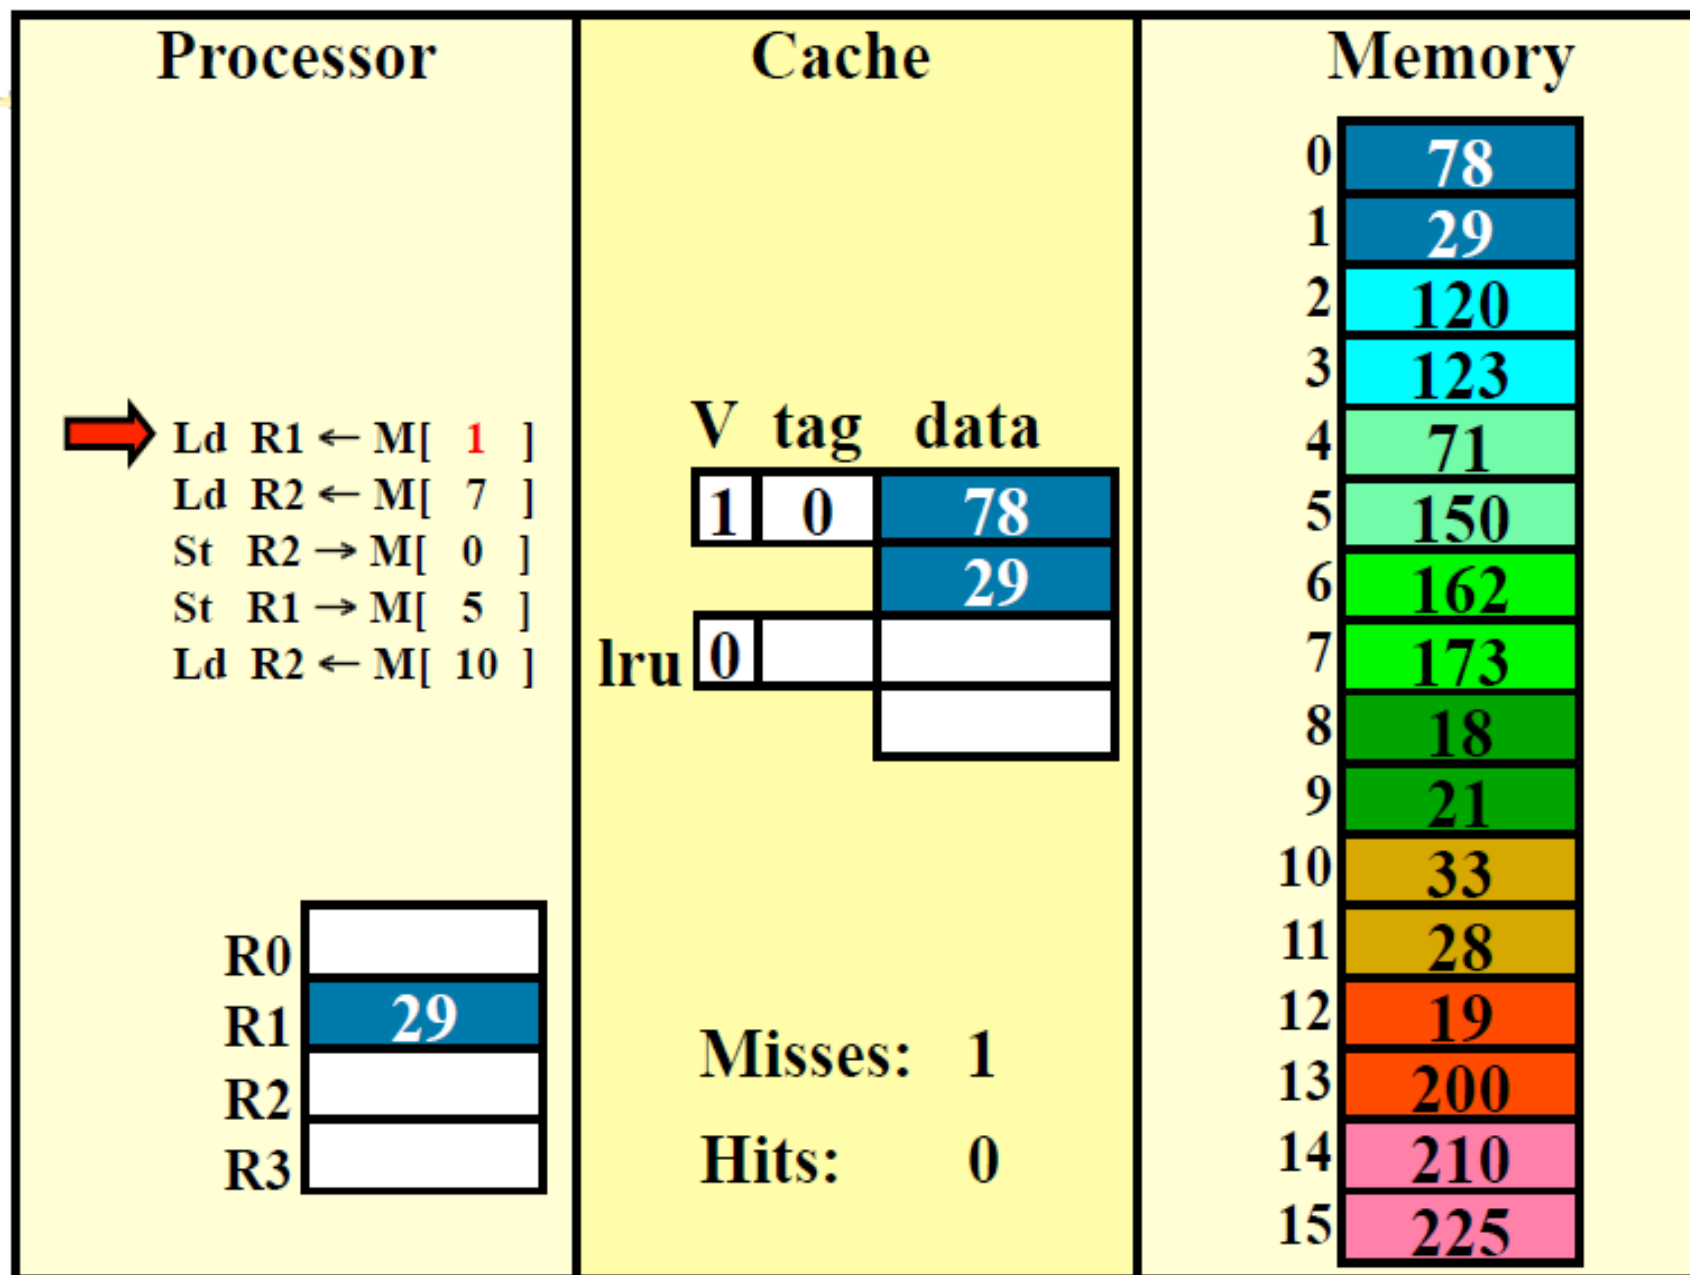
Handling stores (write-through)



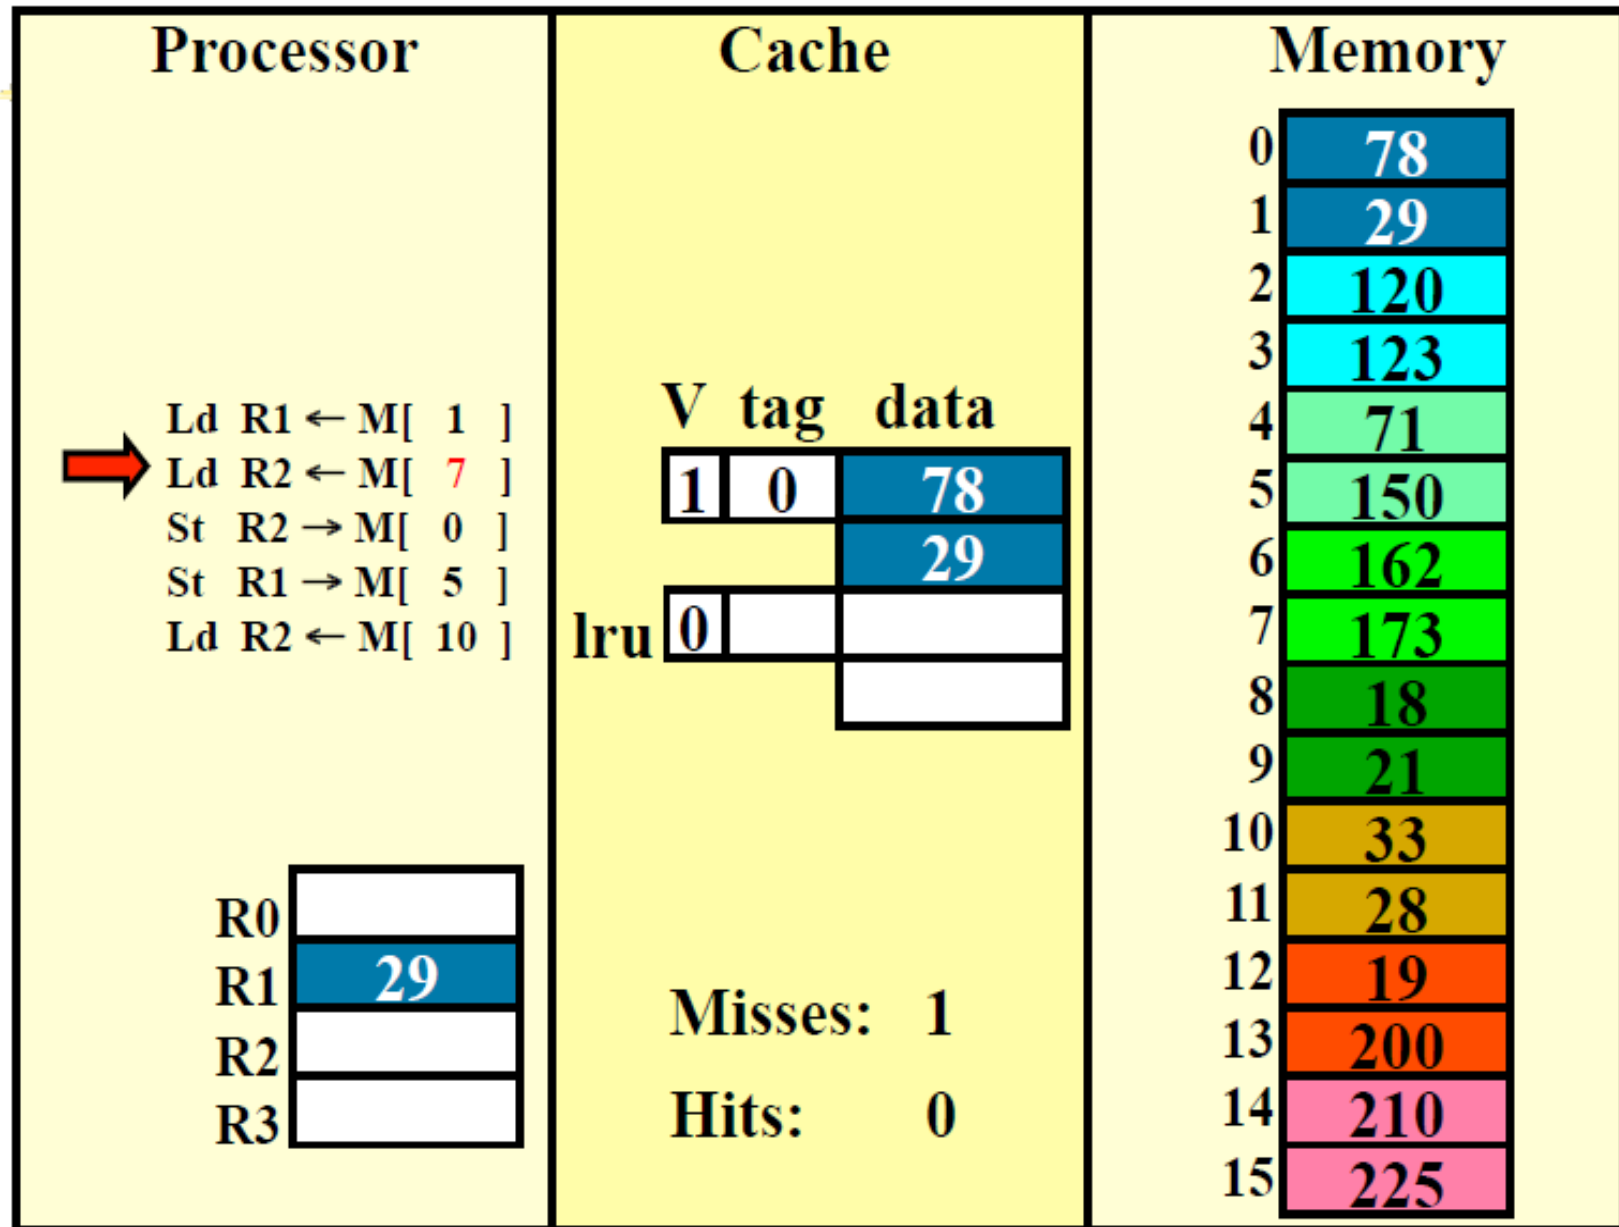
write-through (REF 1)



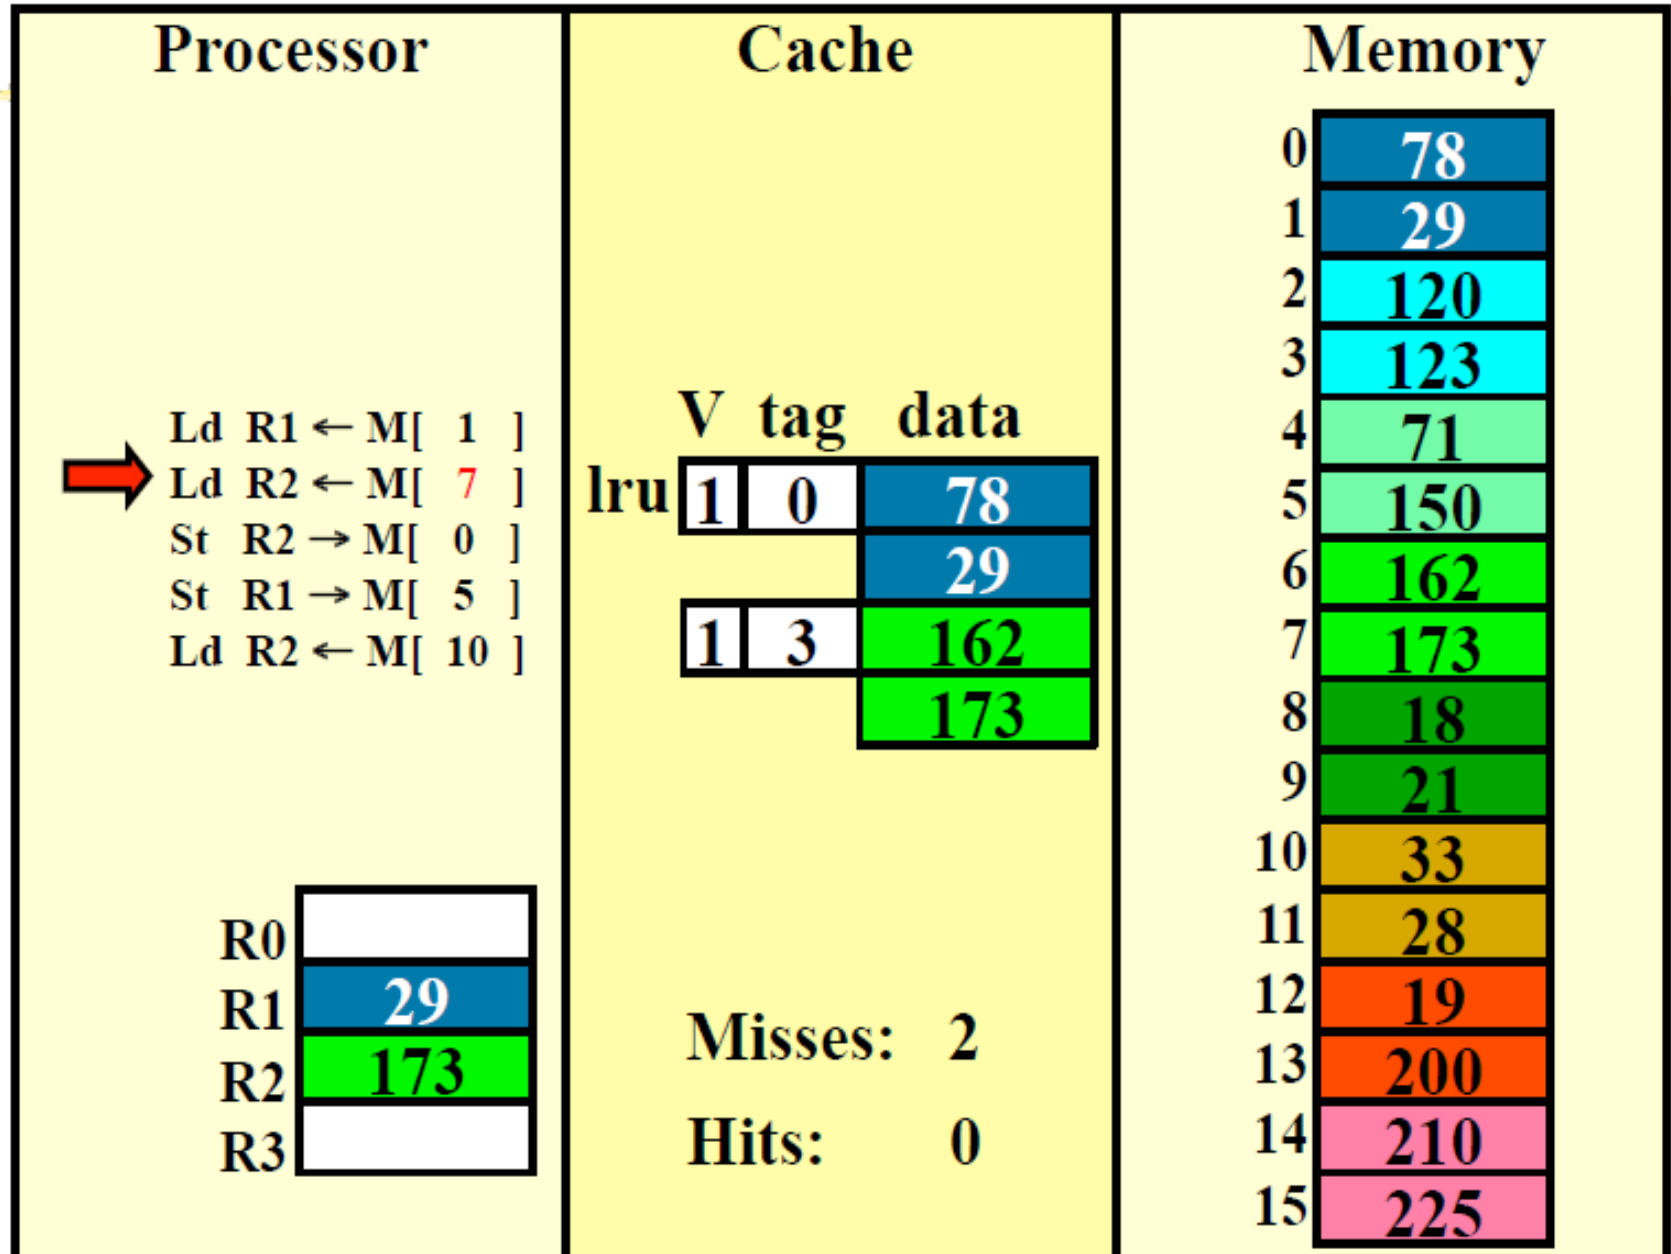
write-through (REF 1)



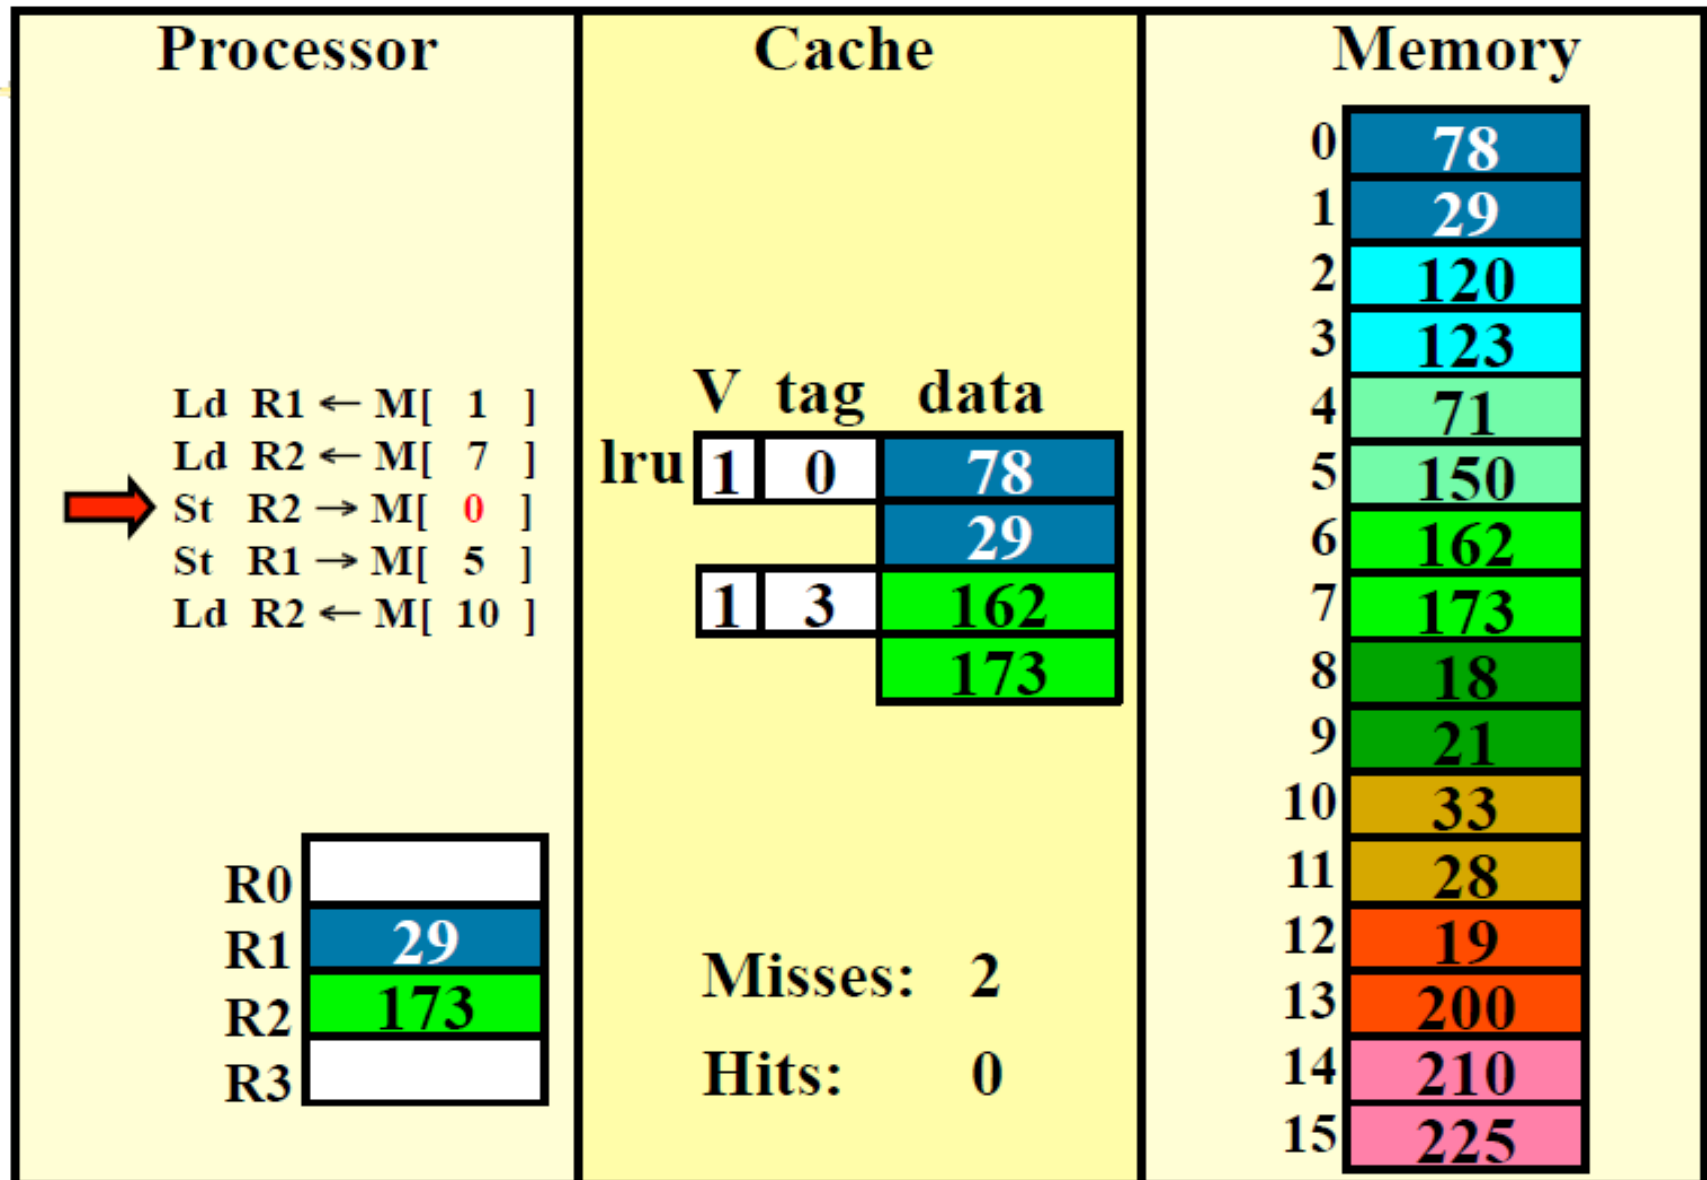
write-through (REF 2)



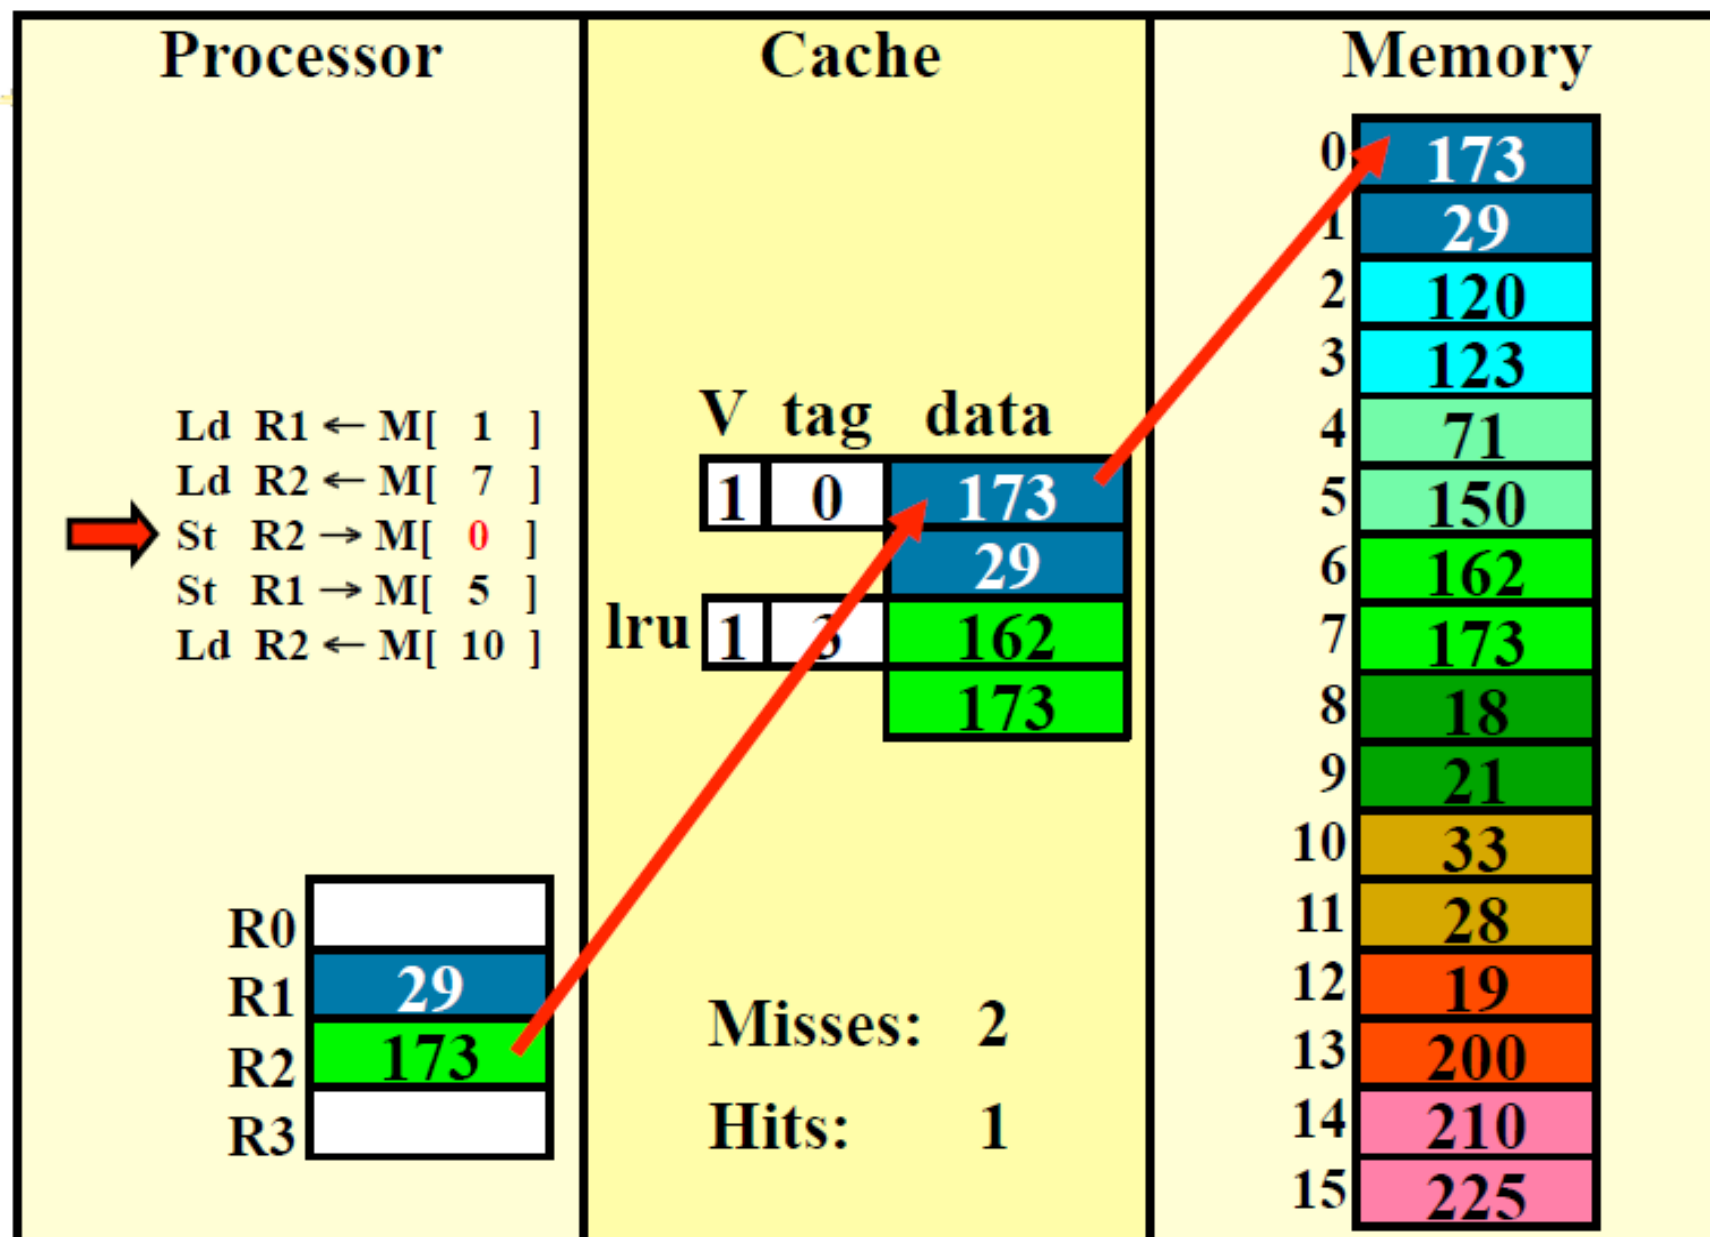
write-through (REF 2)



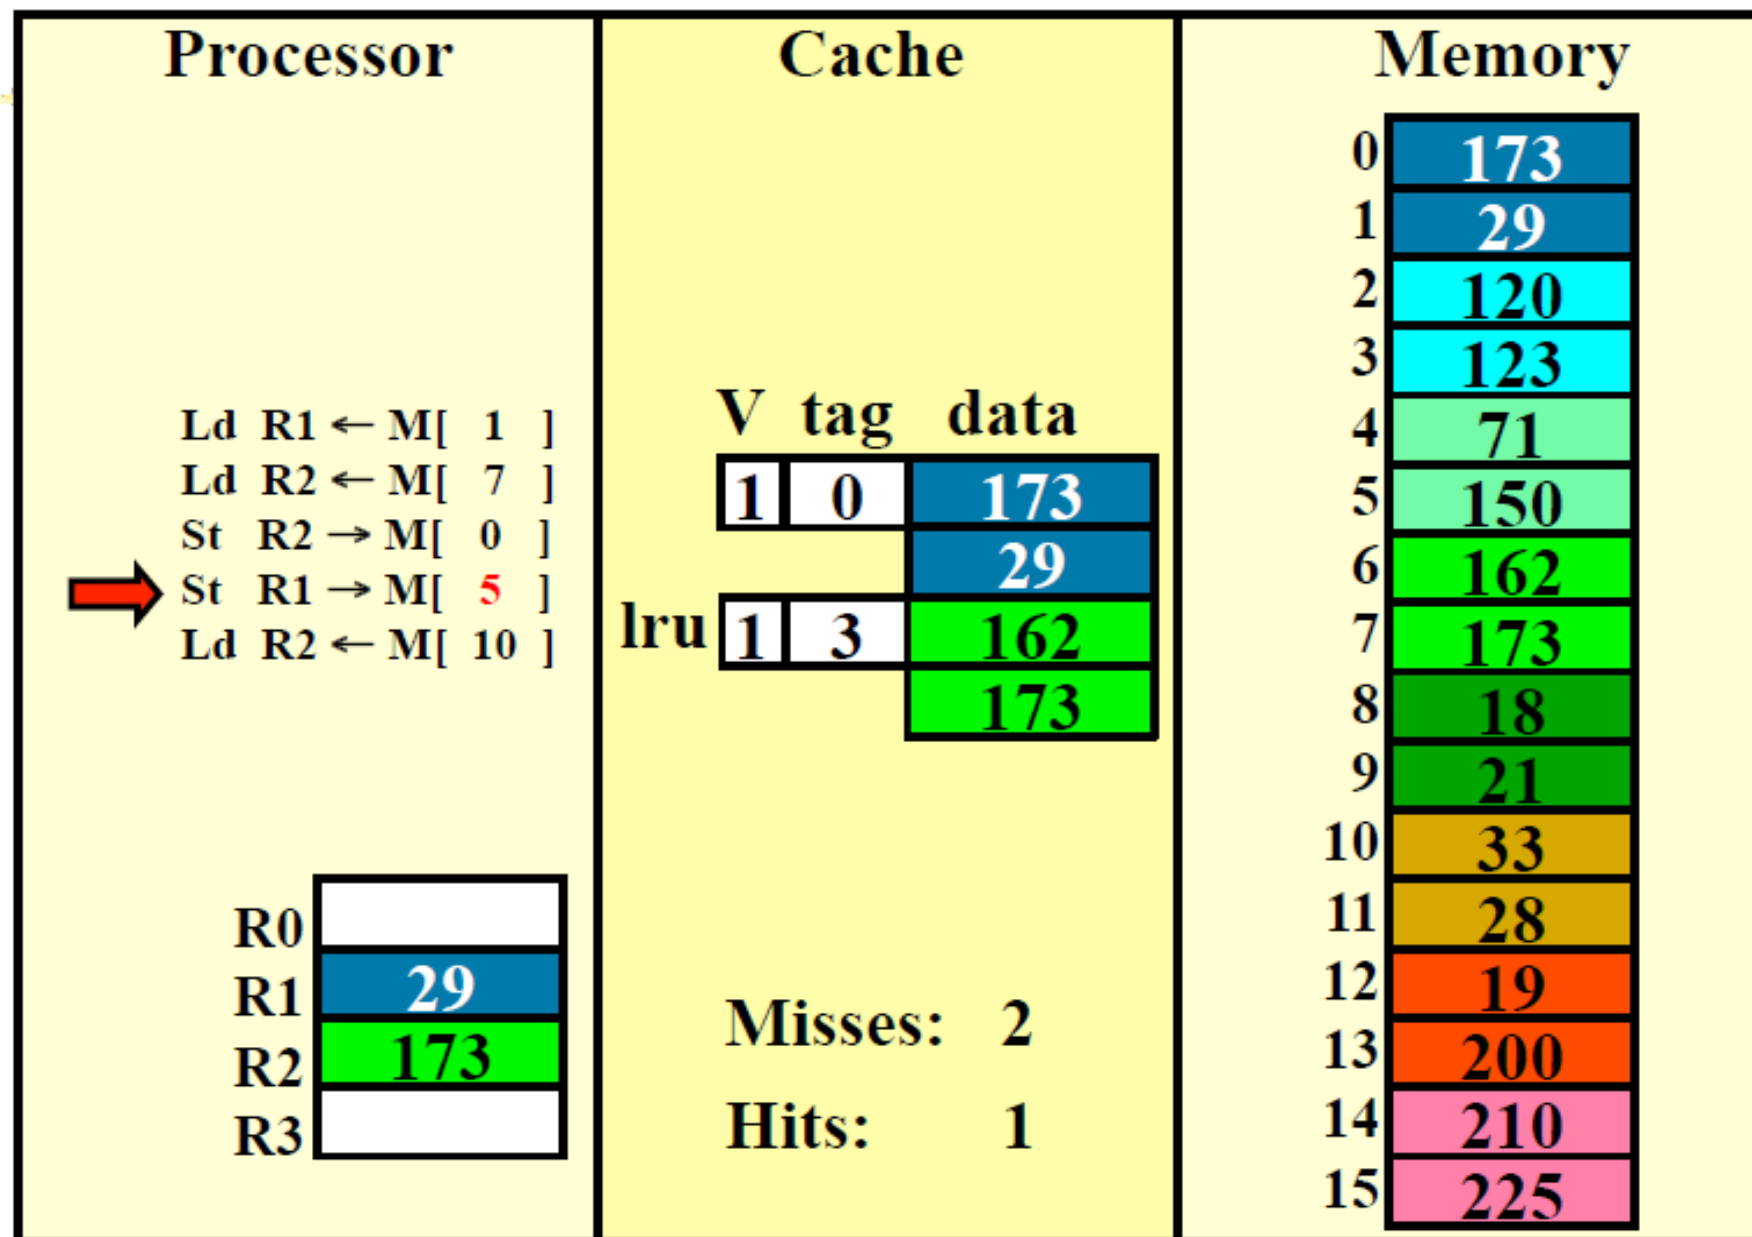
write-through (REF 3)



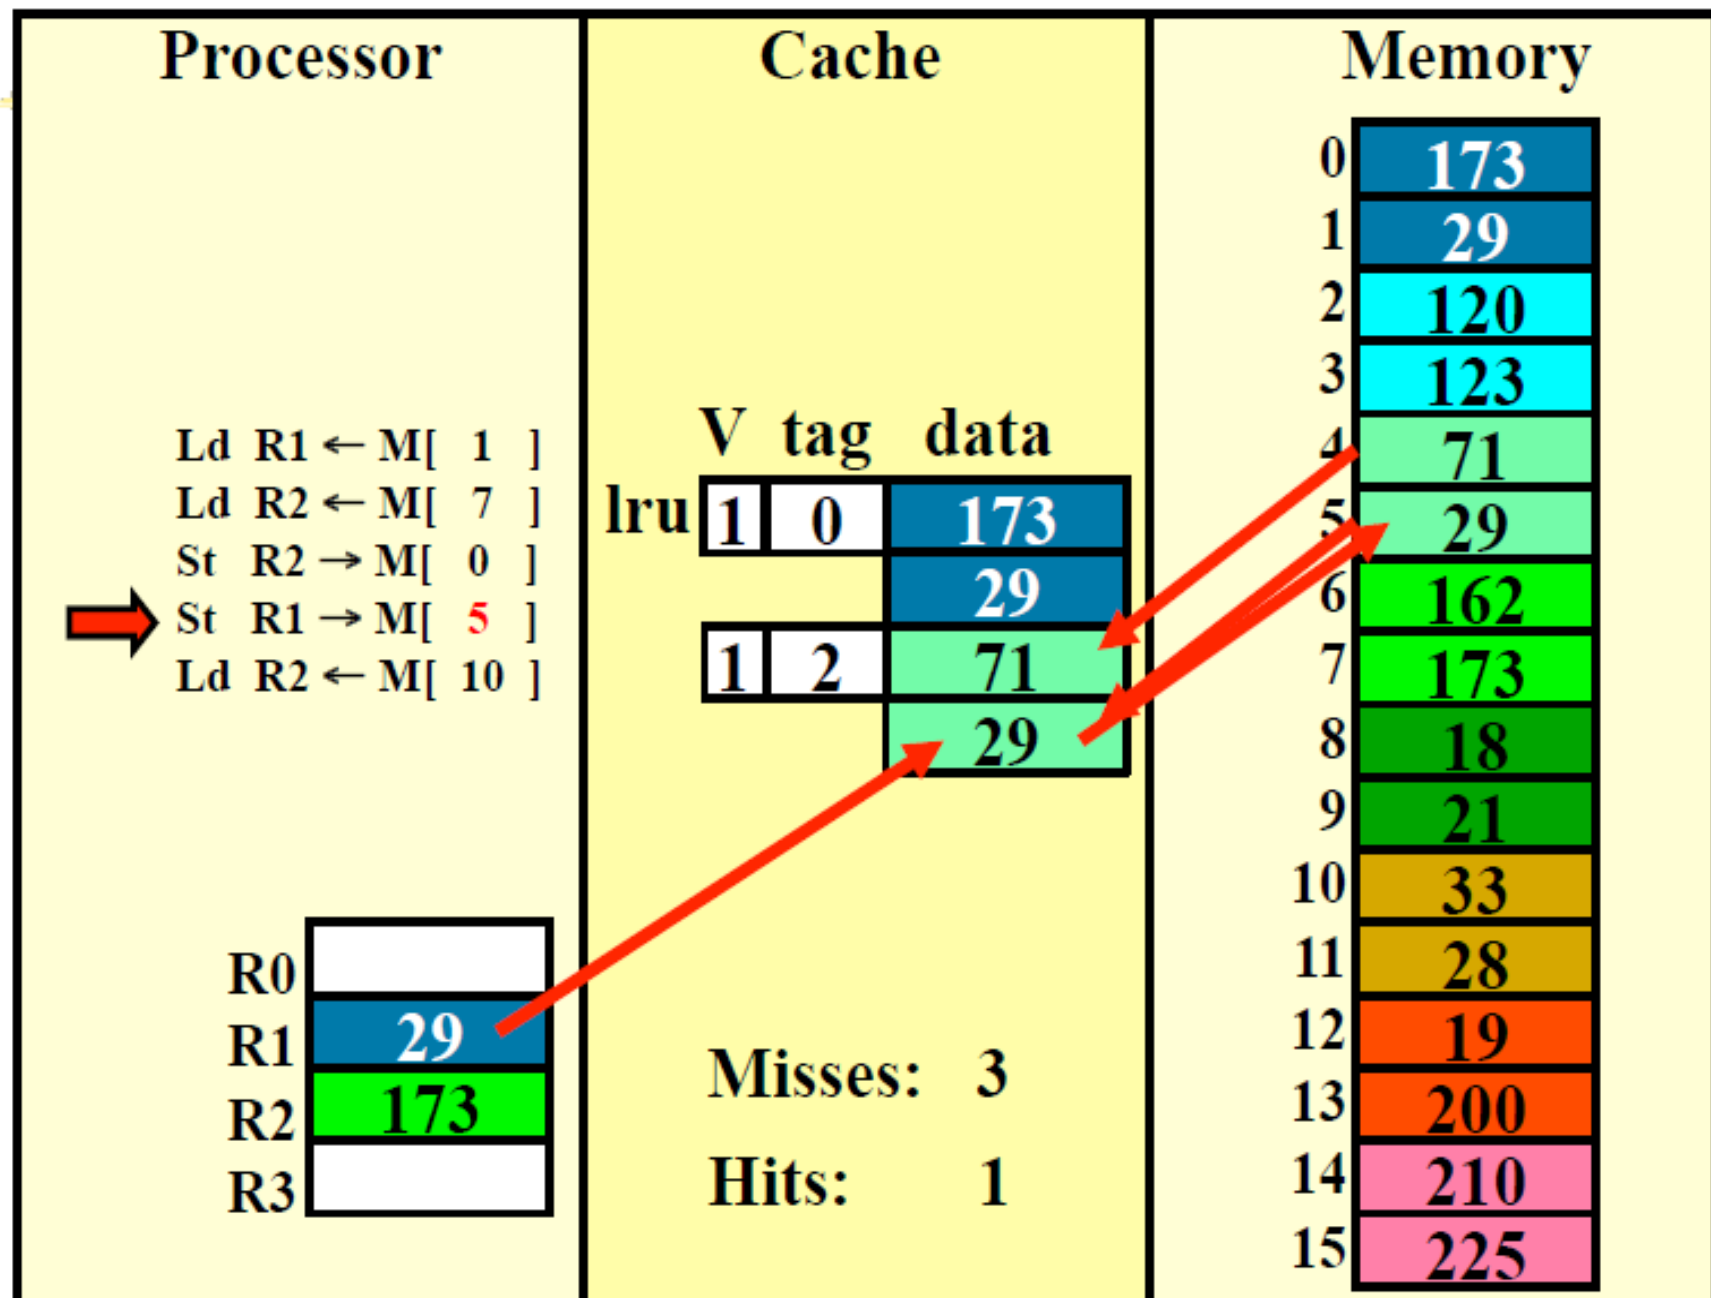
write-through (REF 3)



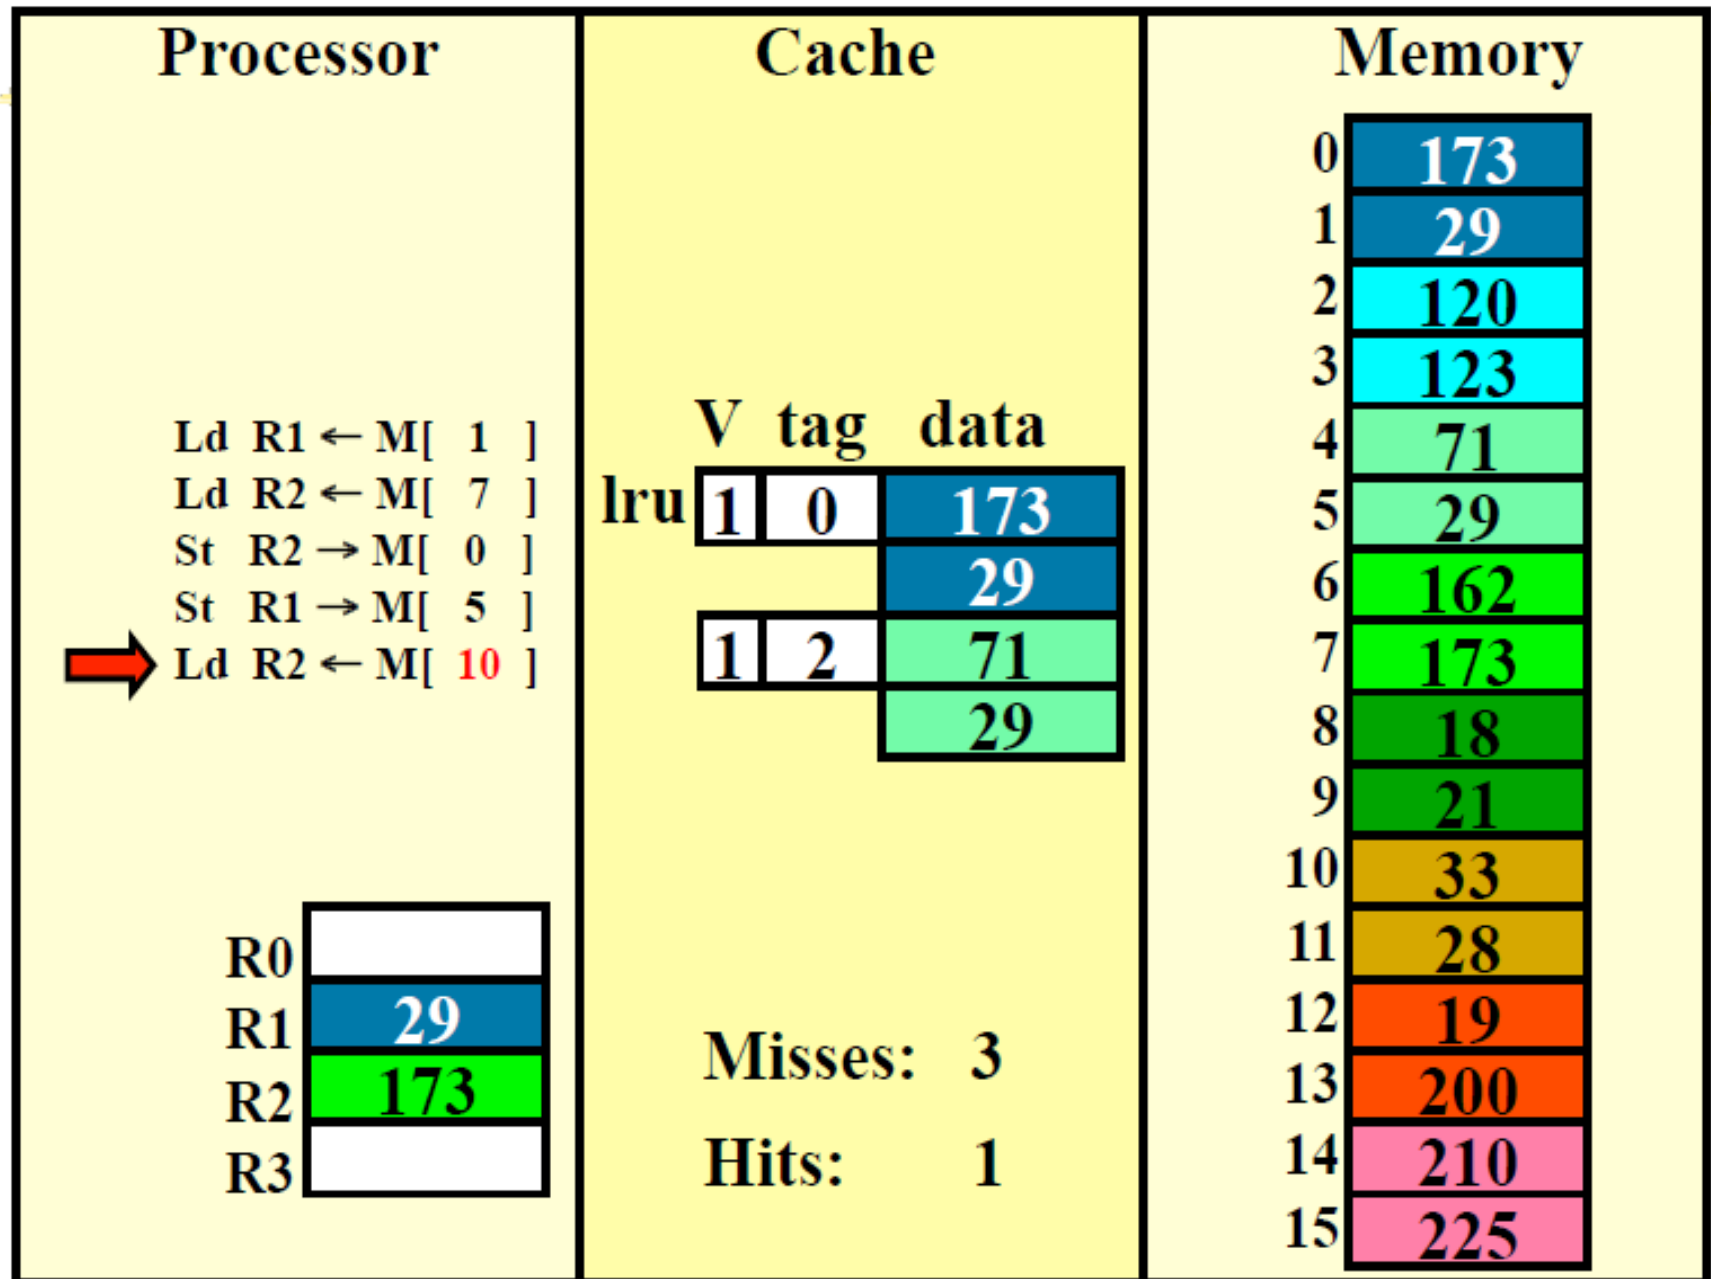
write-through (REF 4)



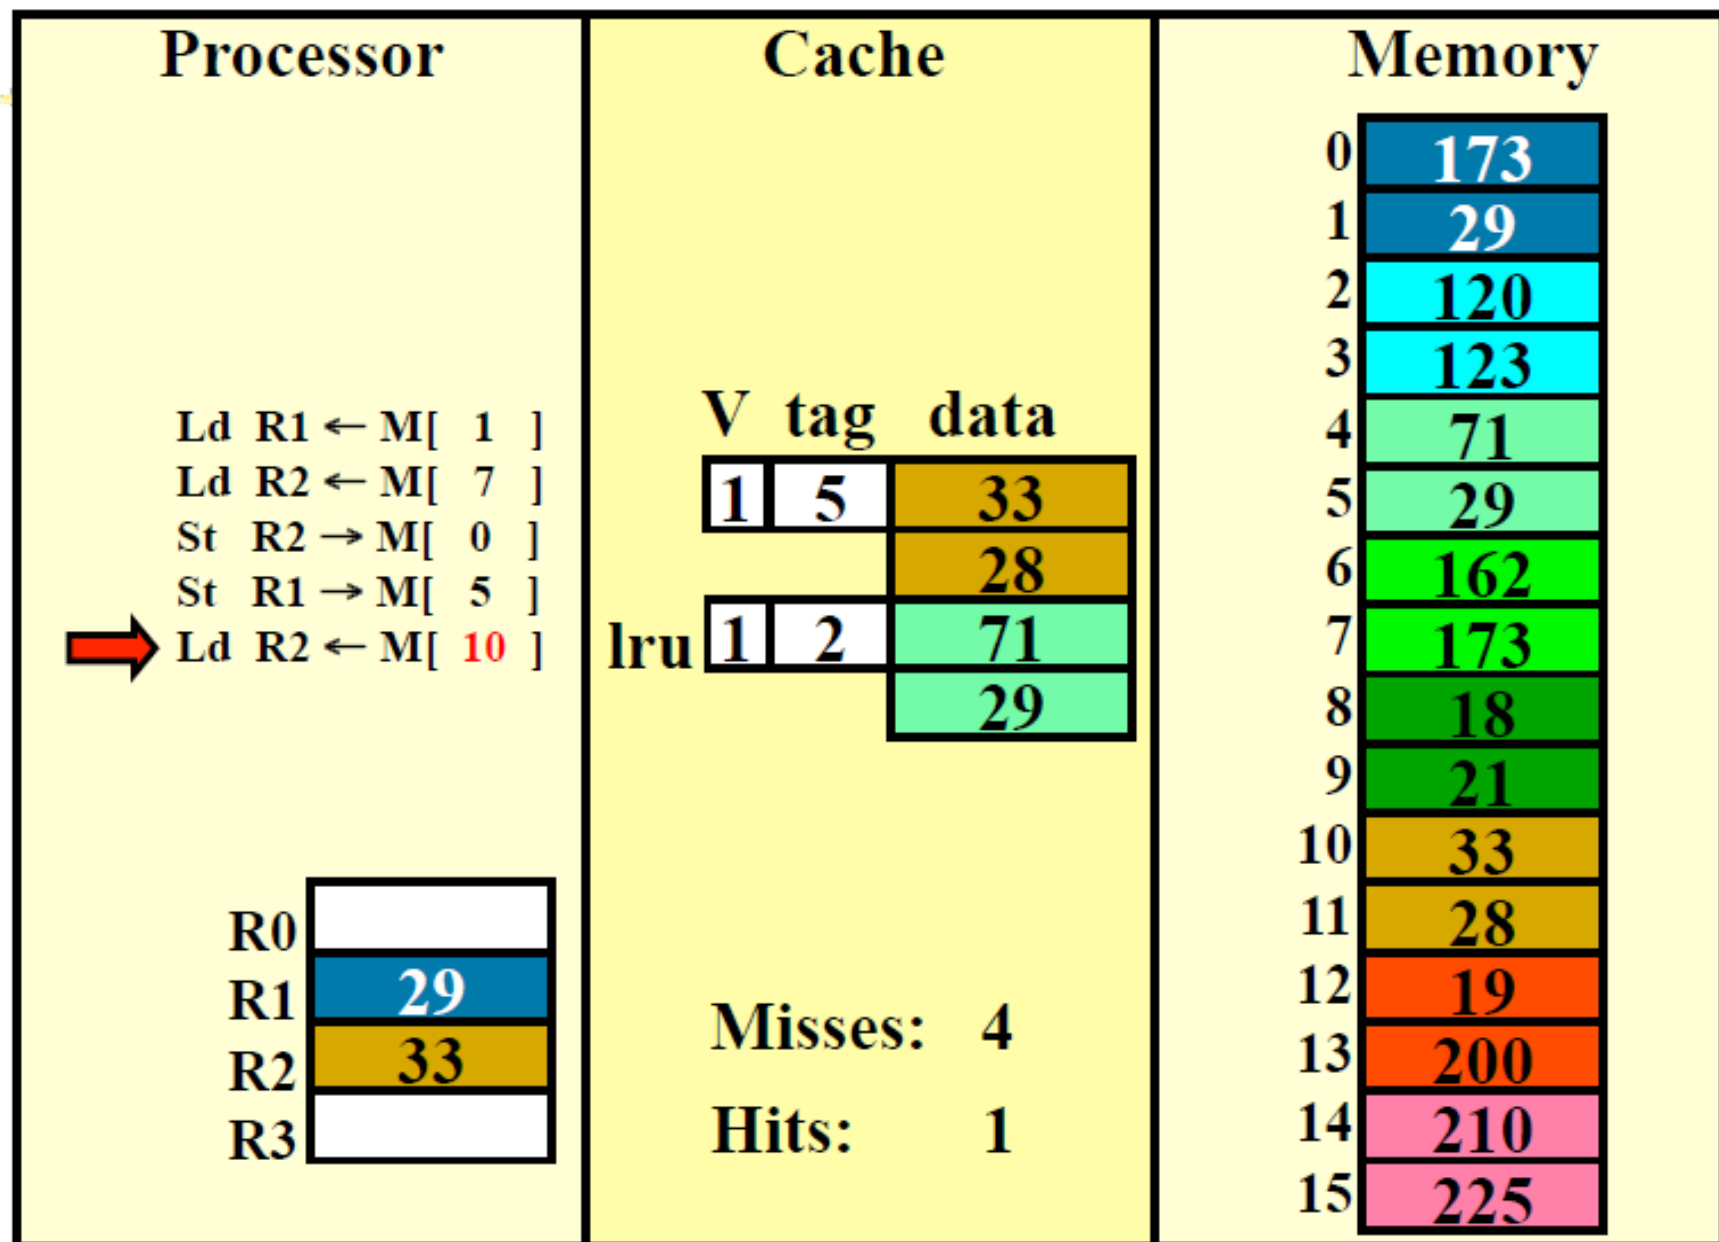
write-through (REF 4)



write-through (REF 5)



write-through (REF 5)

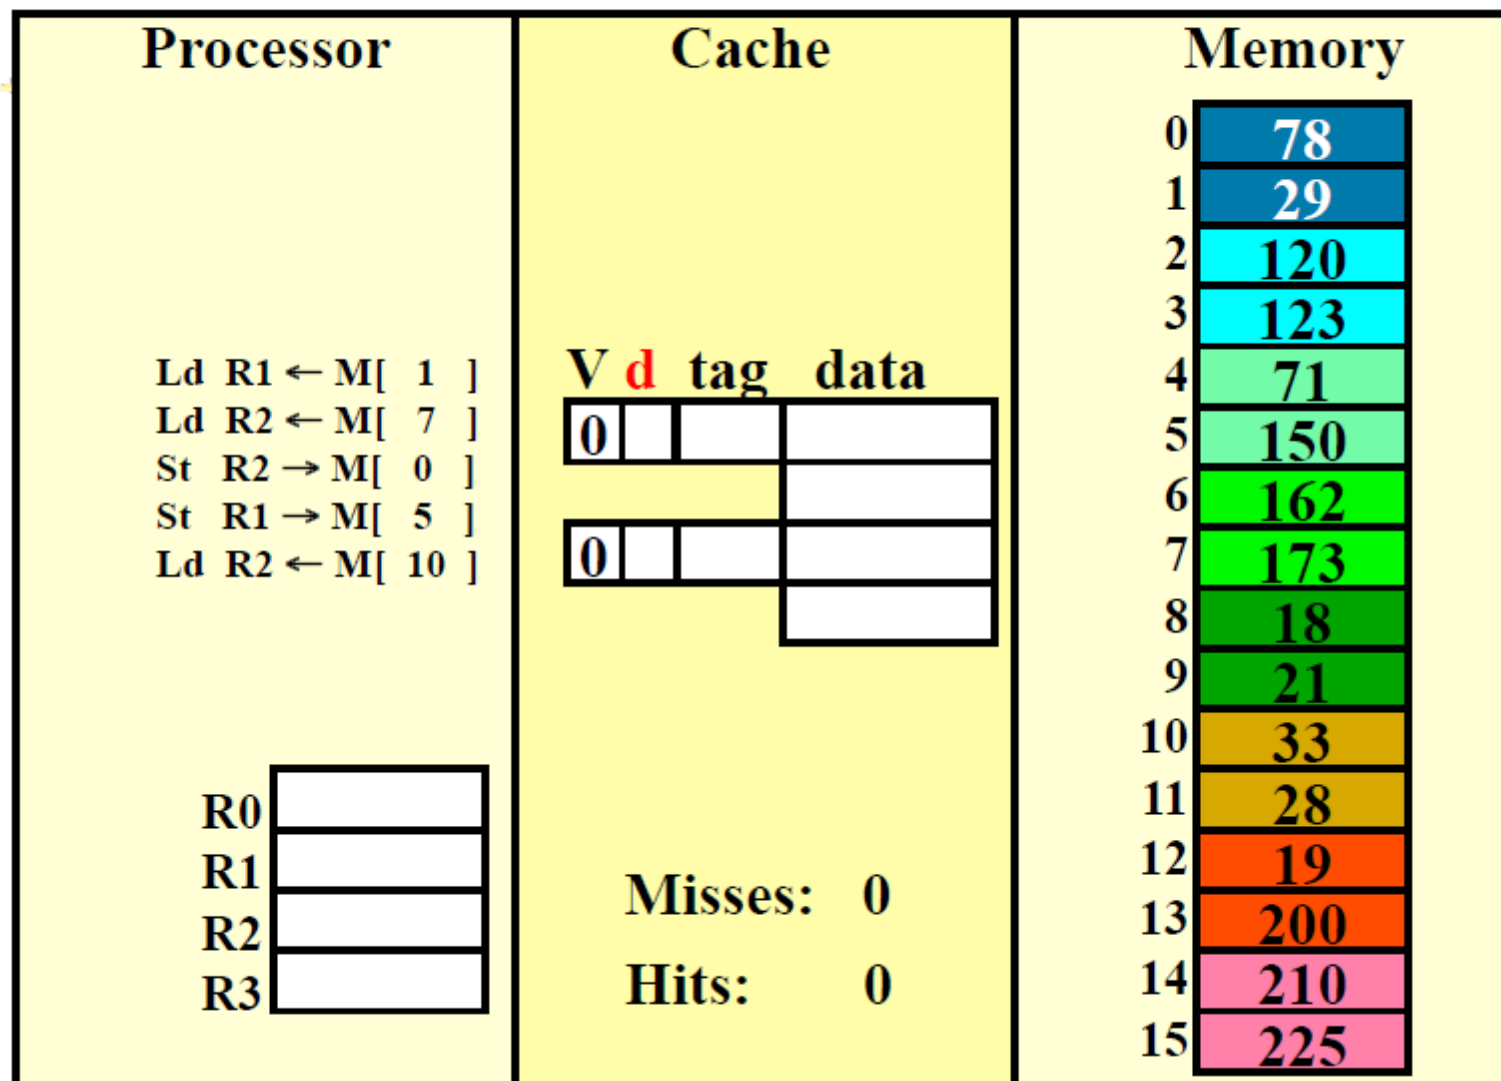


How many memory references?

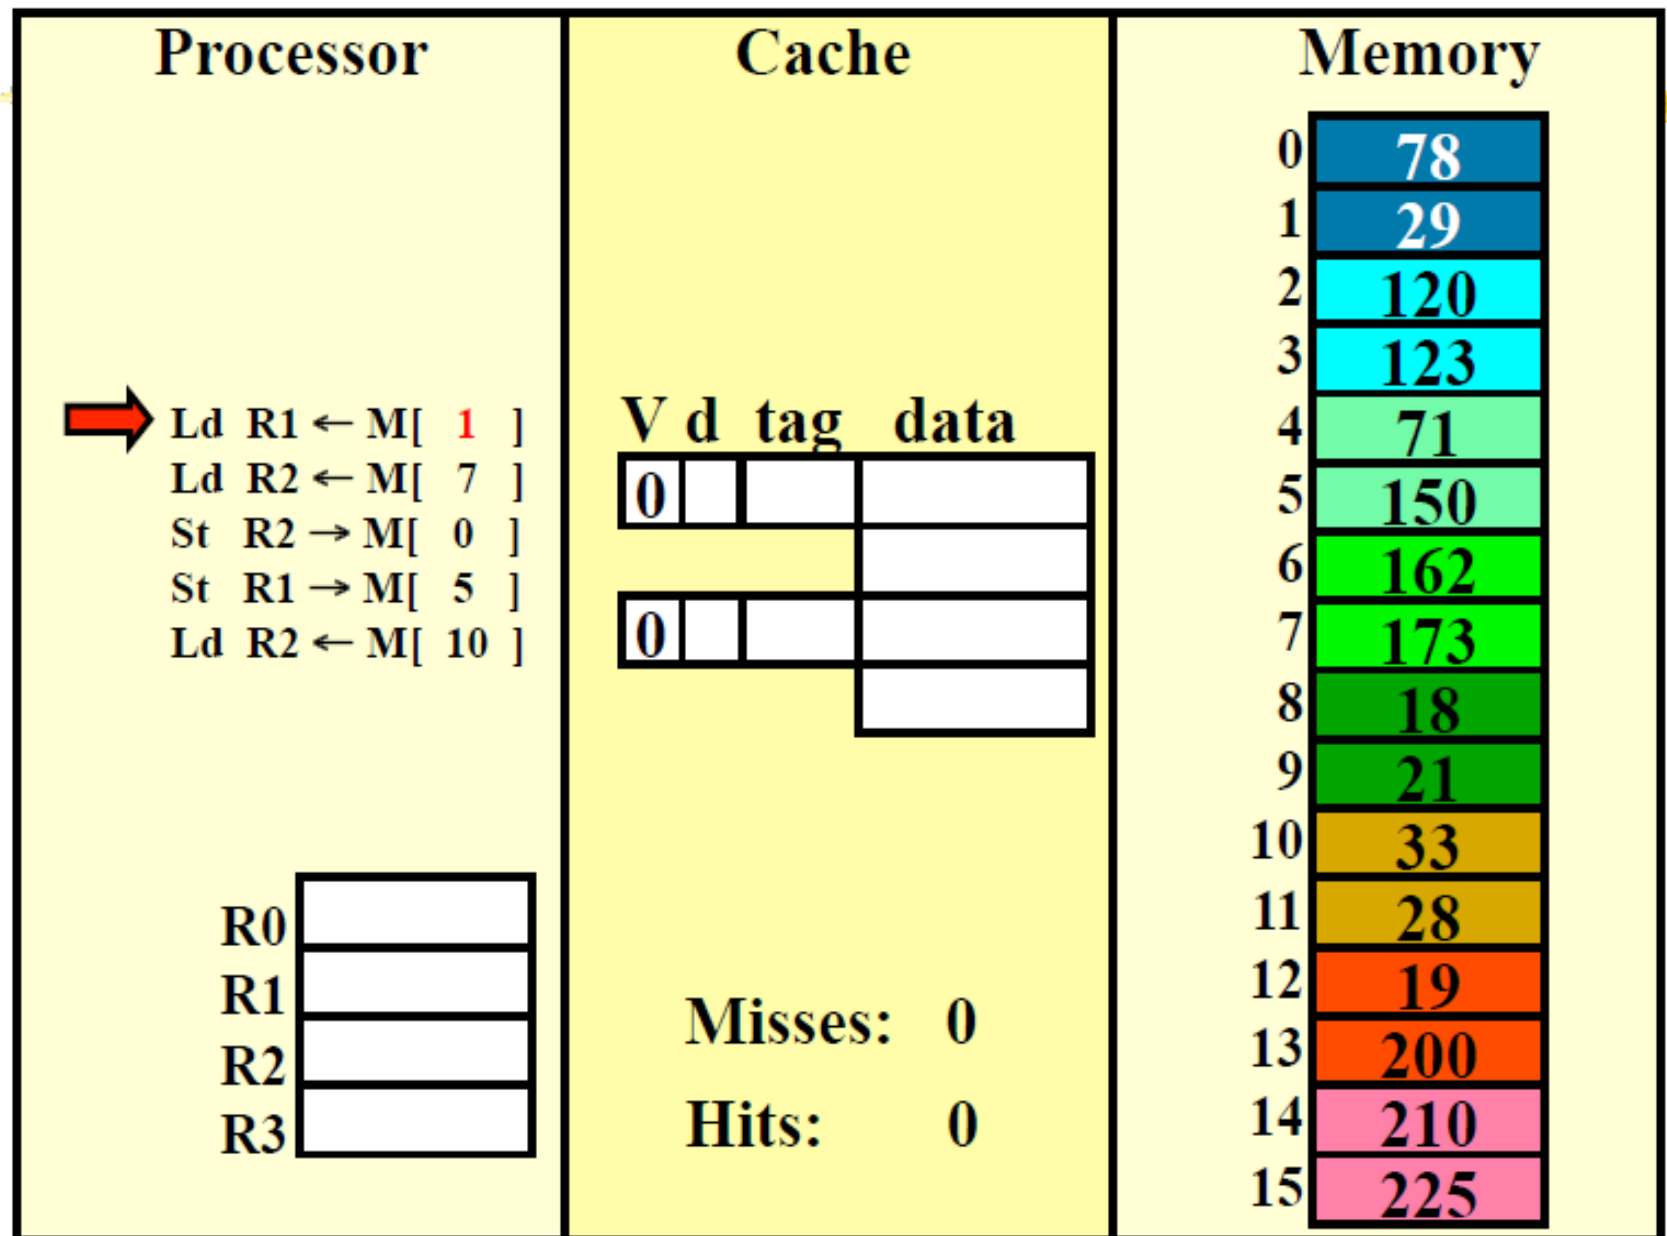
- Each miss reads a block
 - 2 bytes in this cache
- Each store writes a byte
- Total reads: 8 bytes
- Total writes: 2 bytes

but caches generally miss $< 20\%$

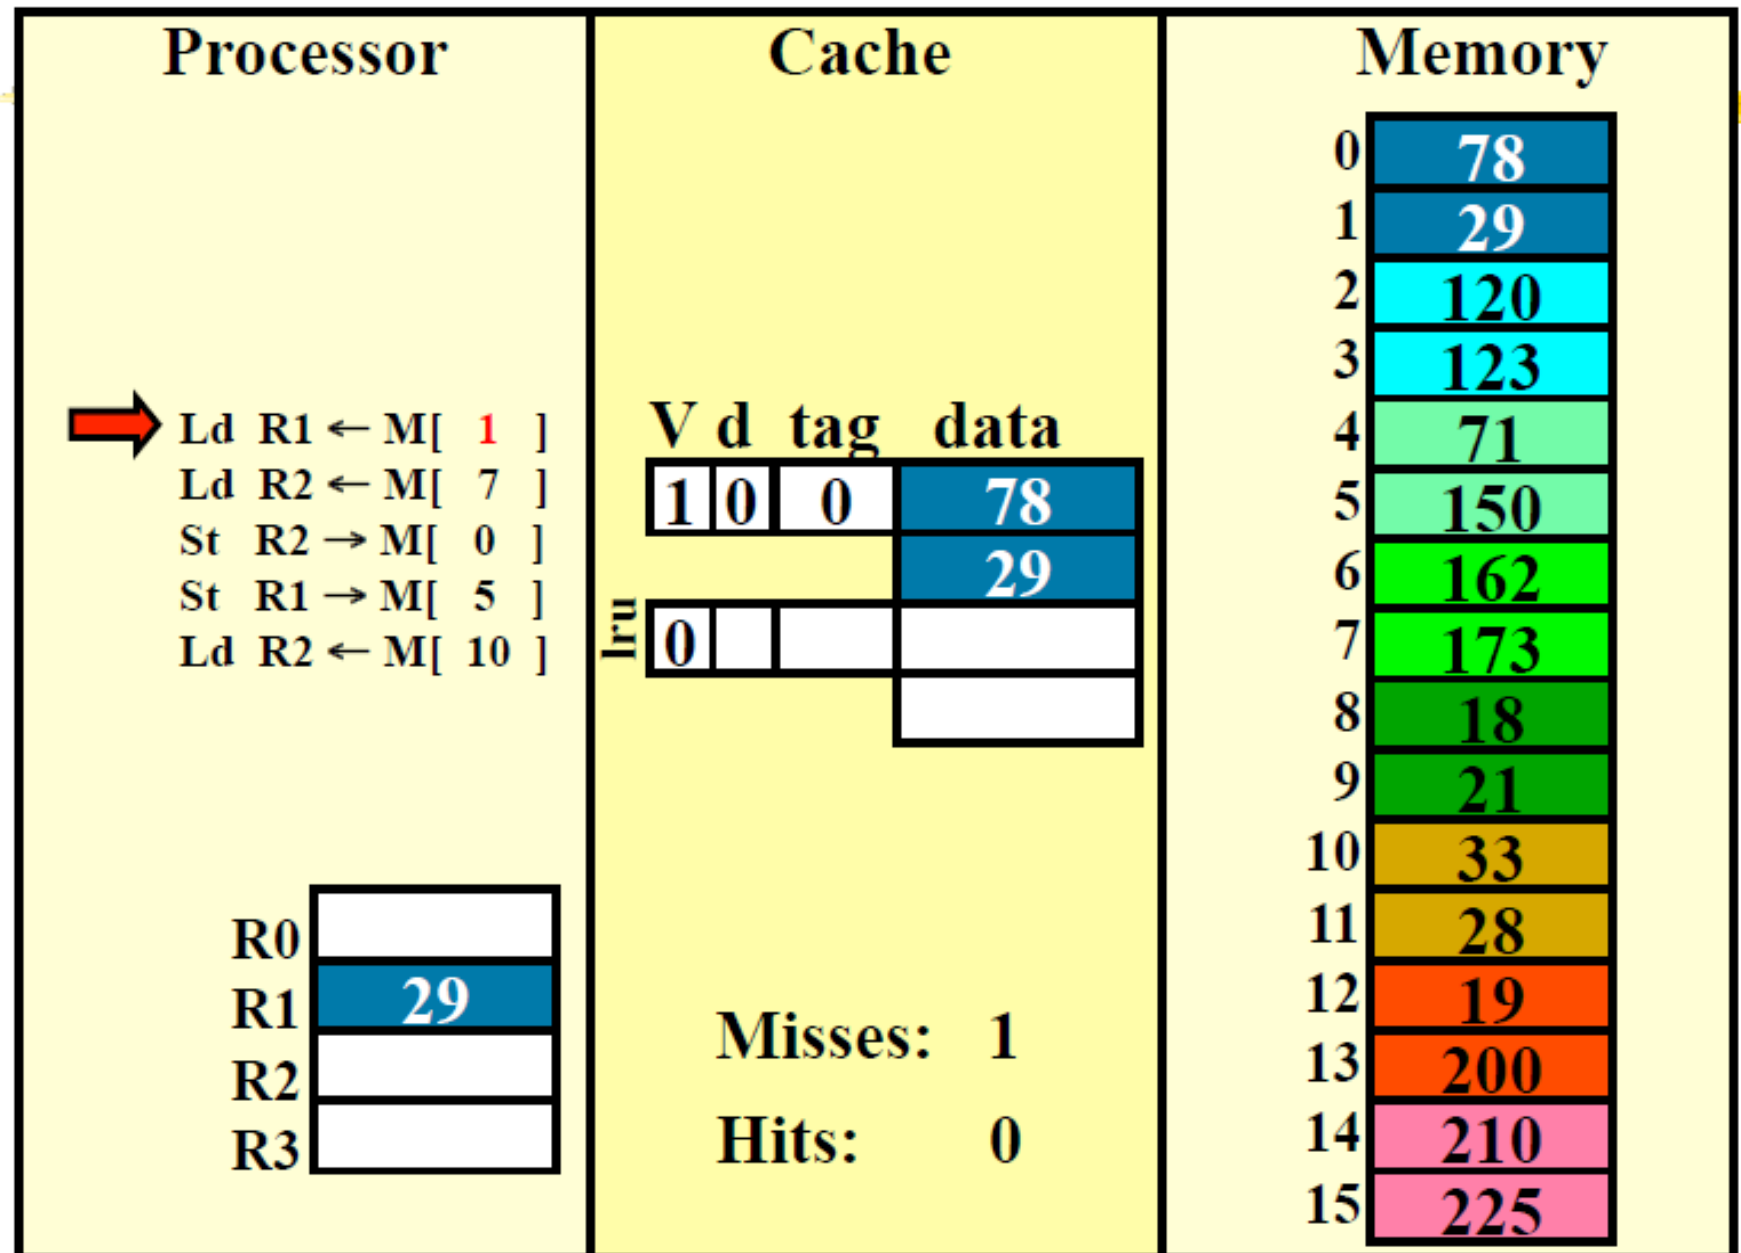
Handling stores (write-back)



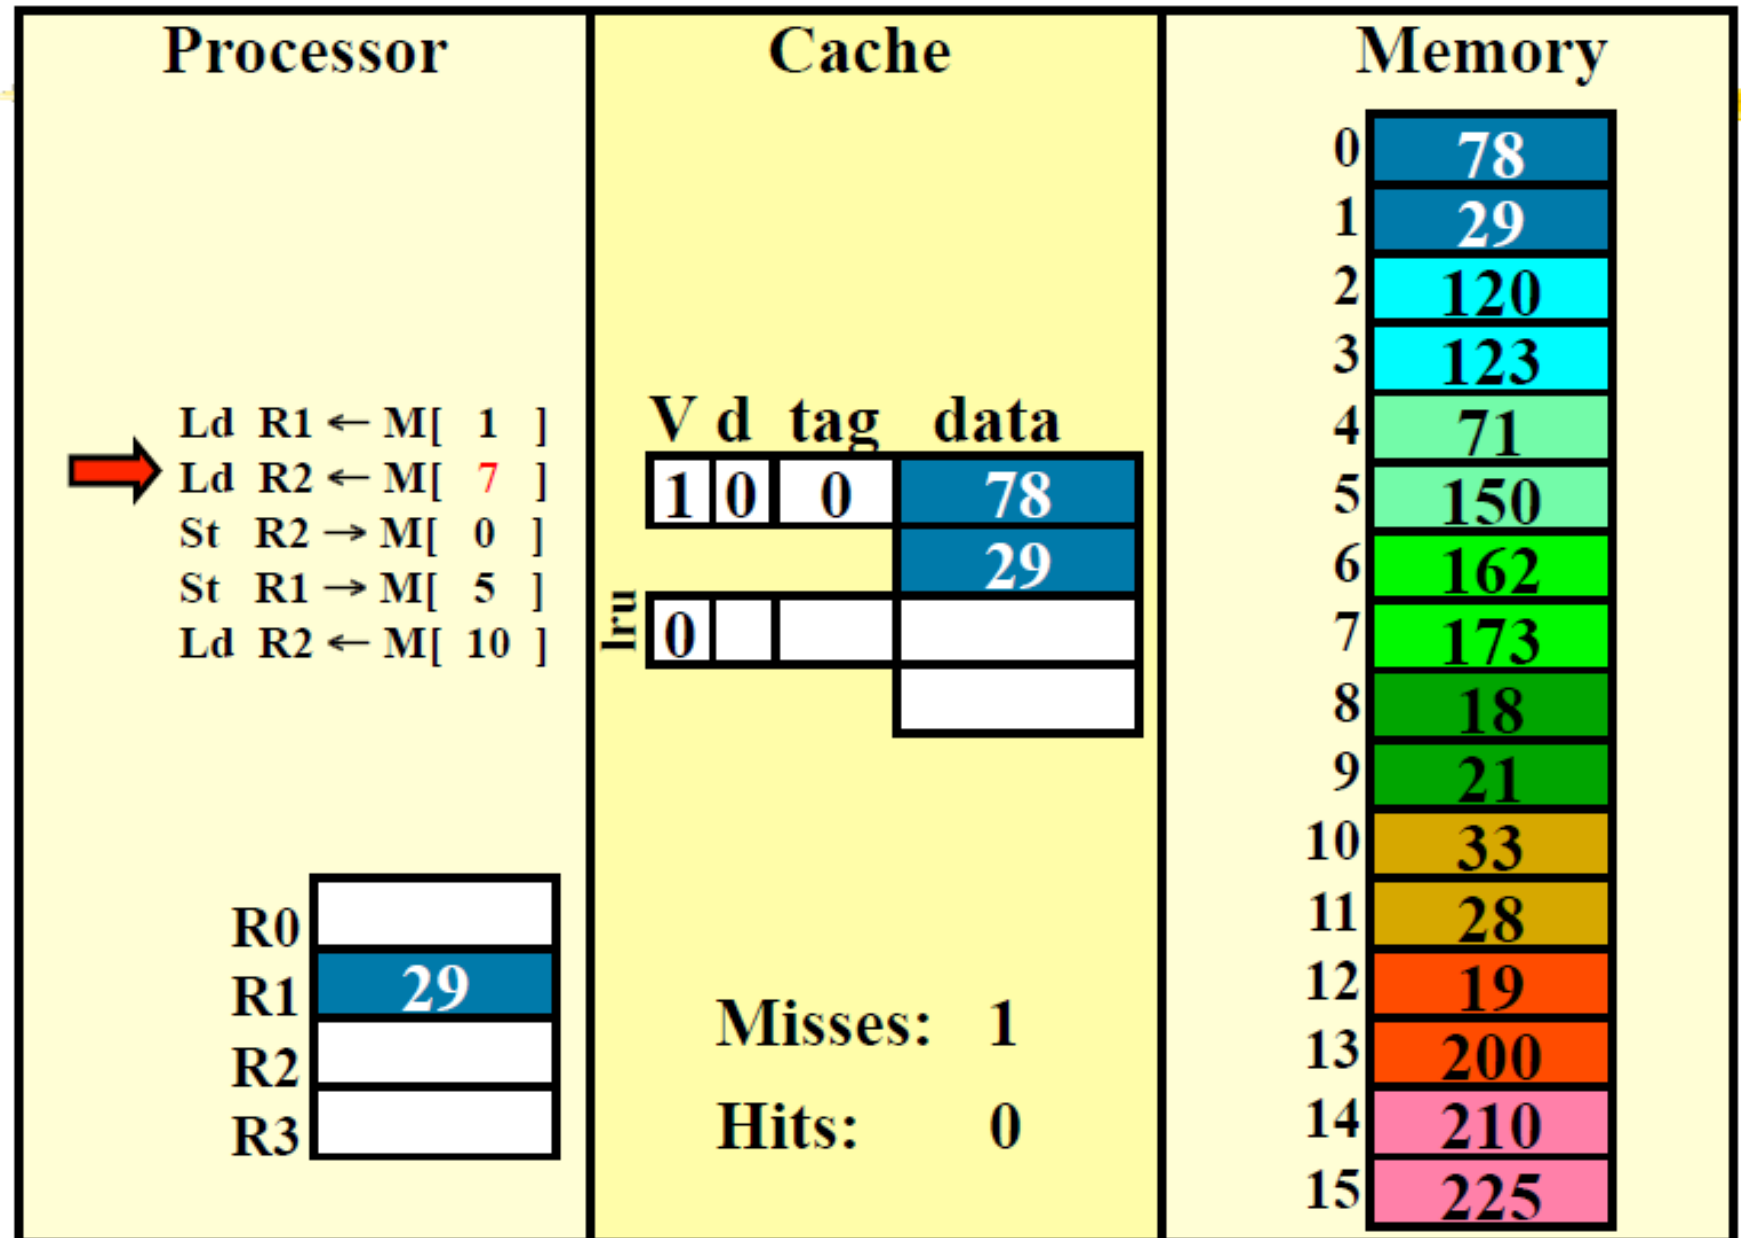
write-back (REF 1)



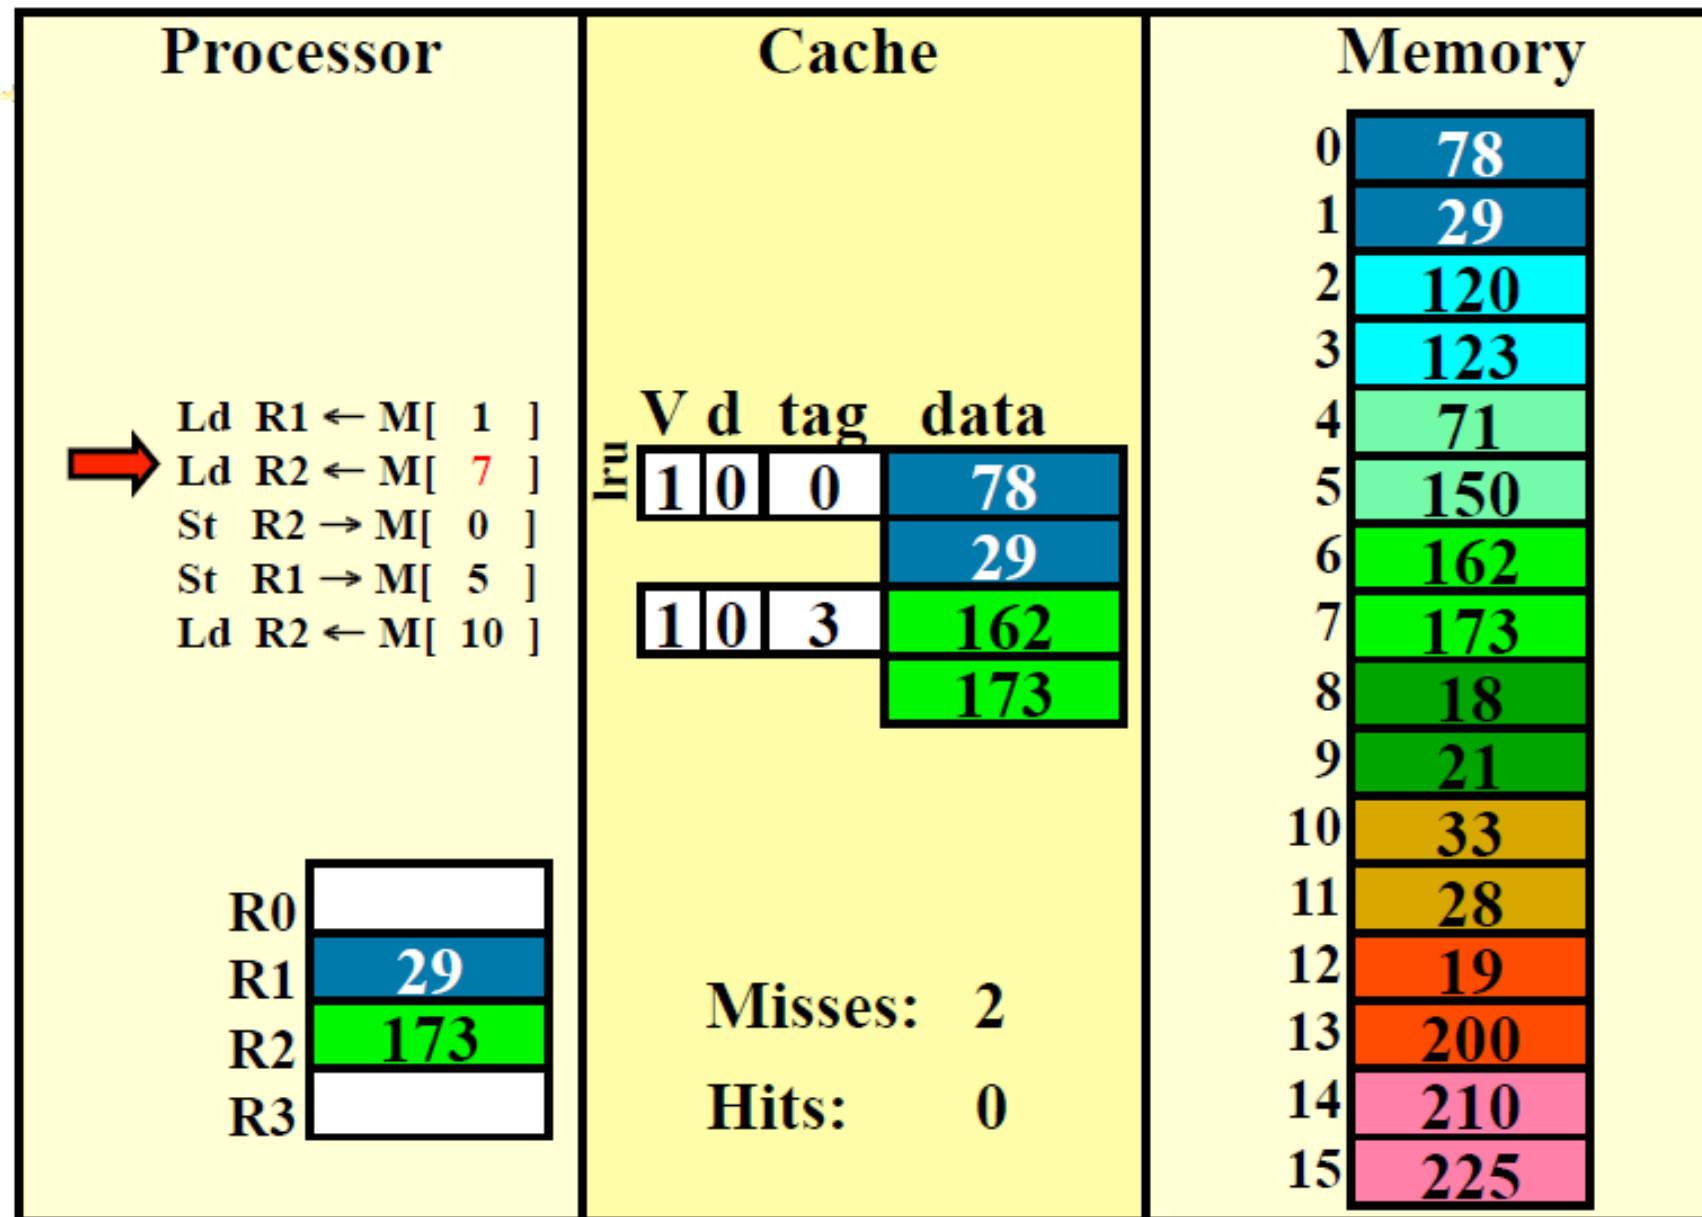
write-back (REF 1)



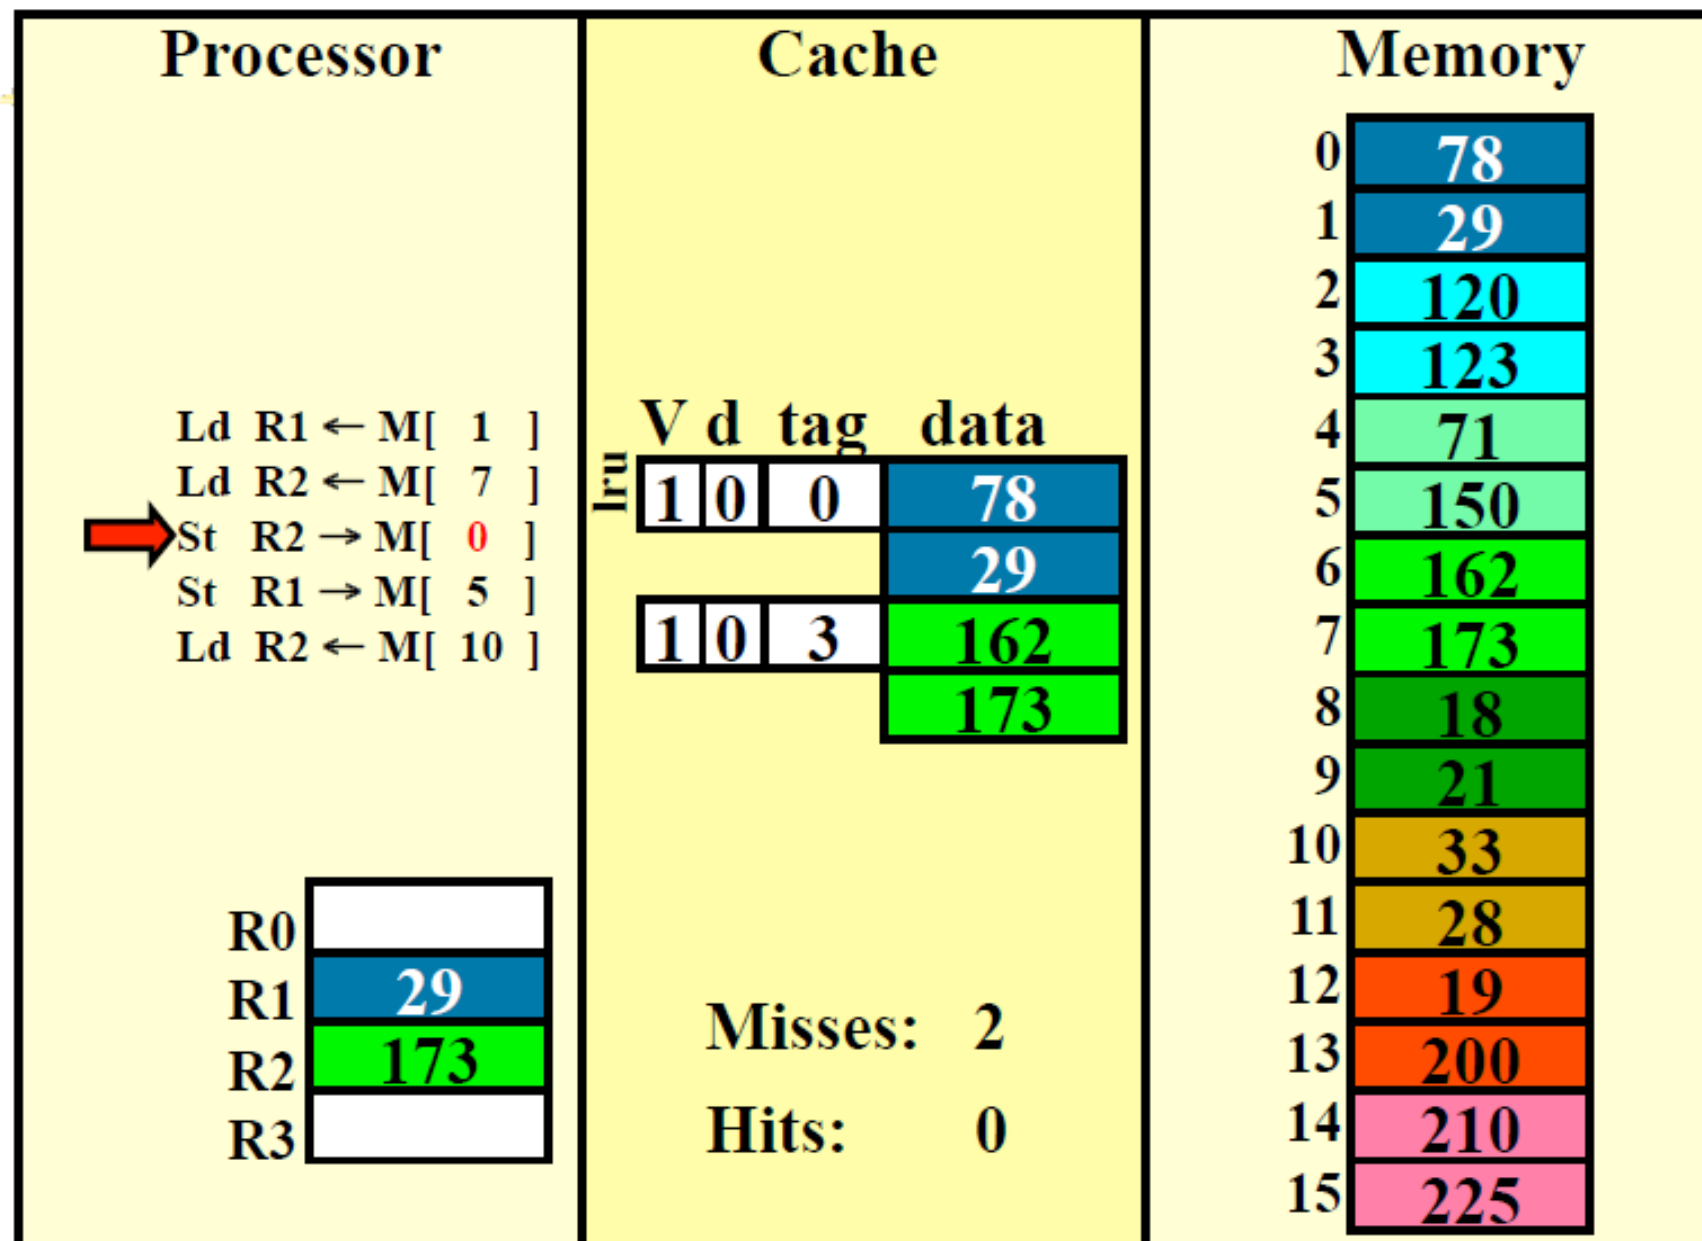
write-back (REF 2)



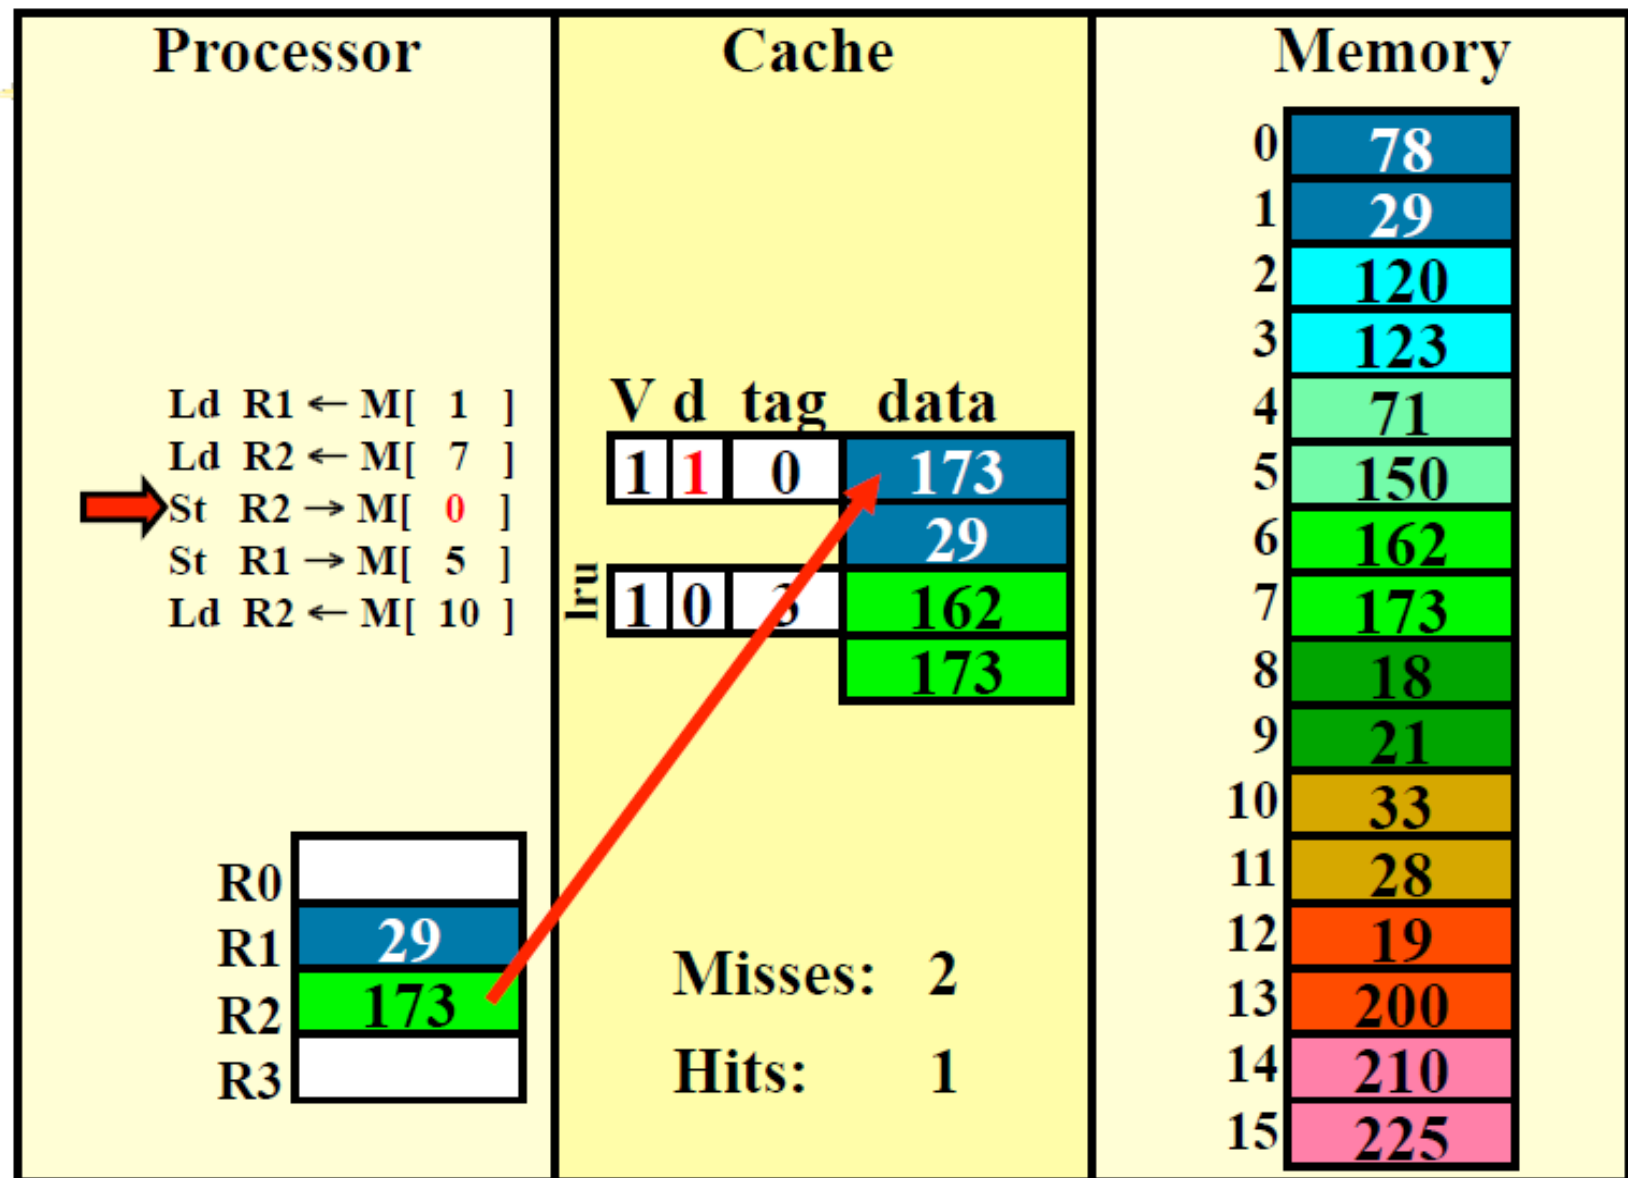
write-back (REF 2)



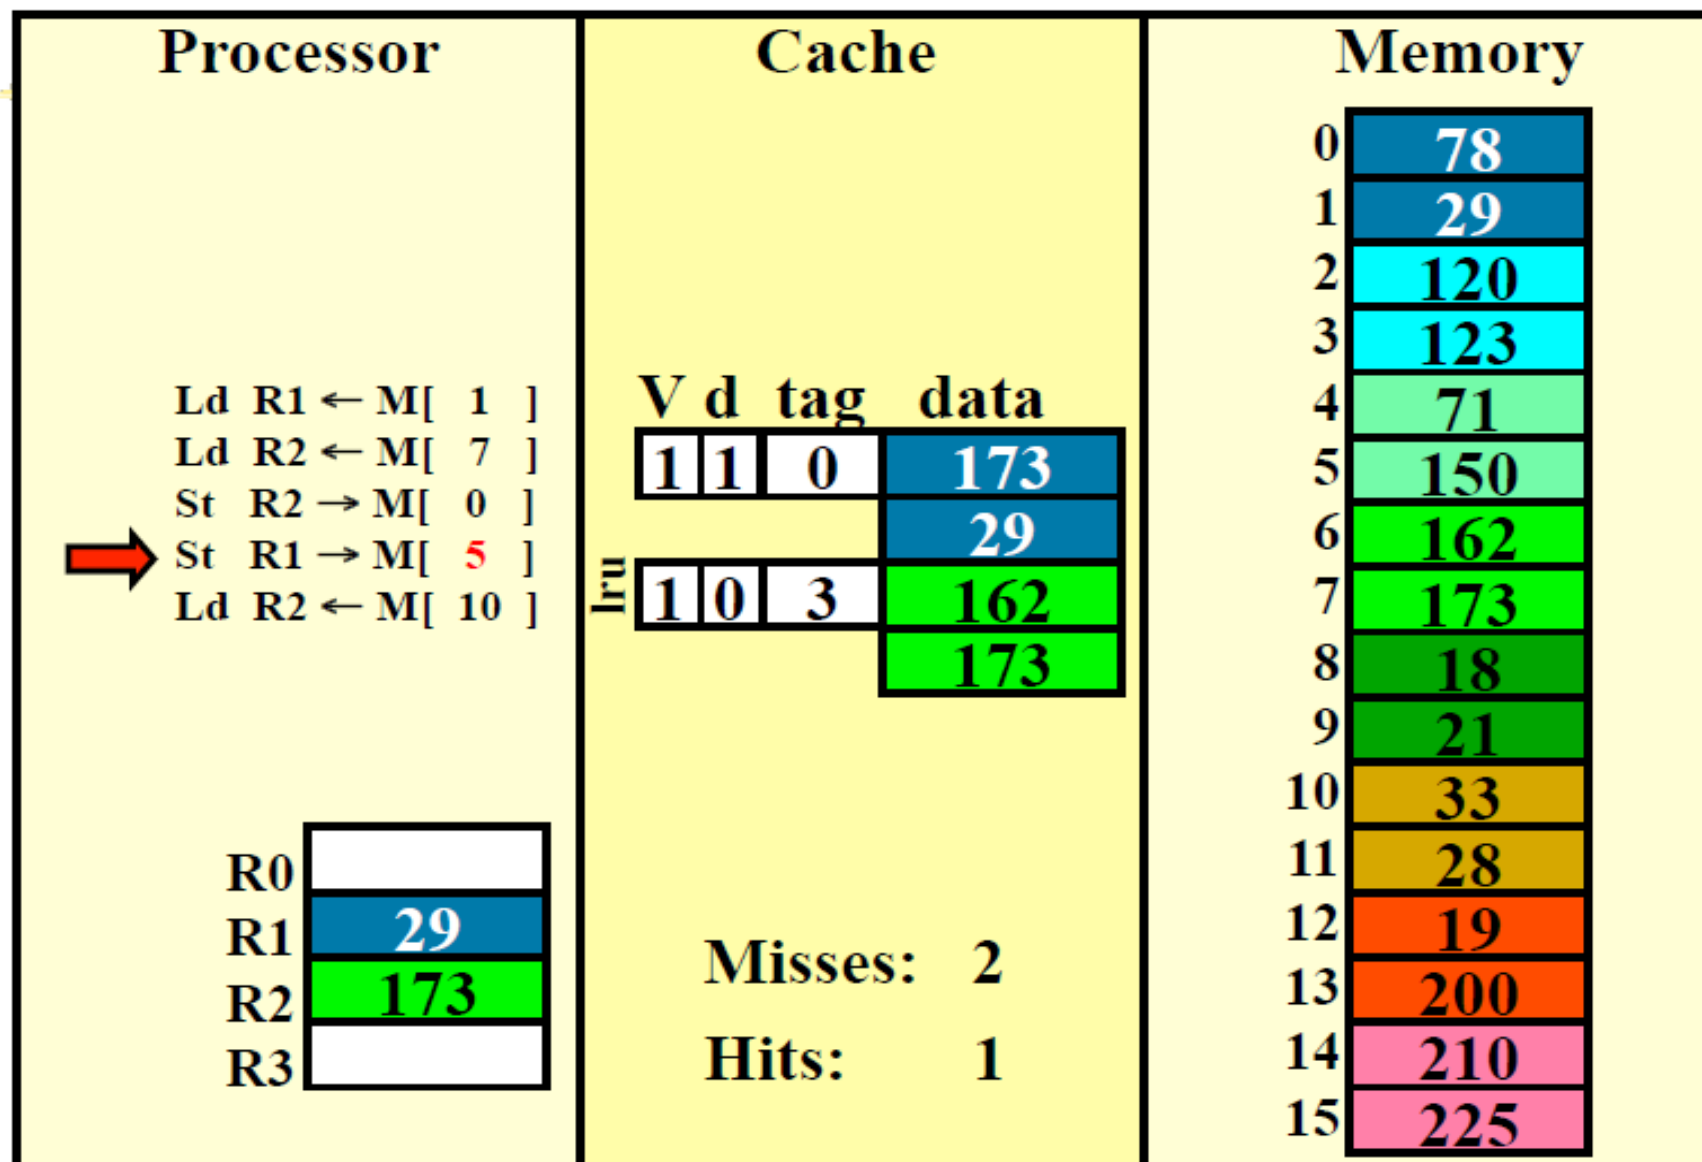
write-back (REF 3)



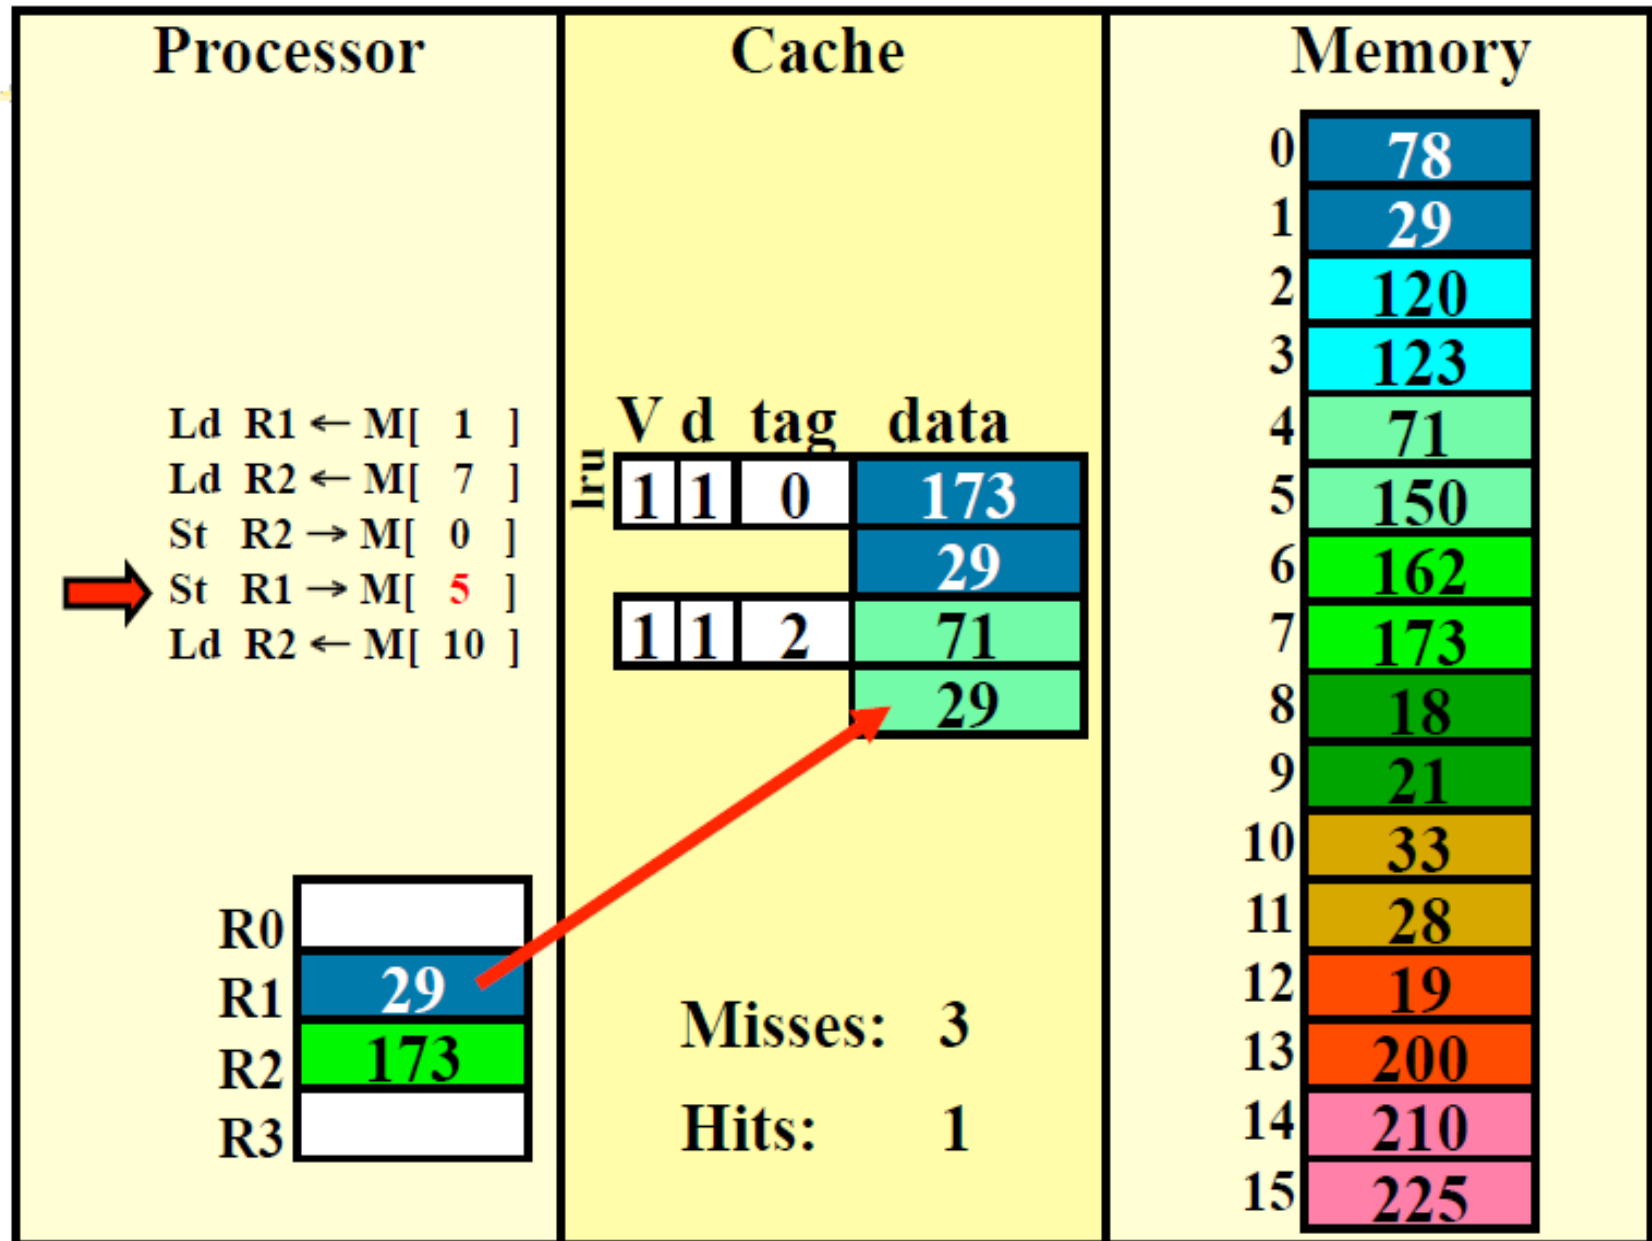
write-back (REF 3)



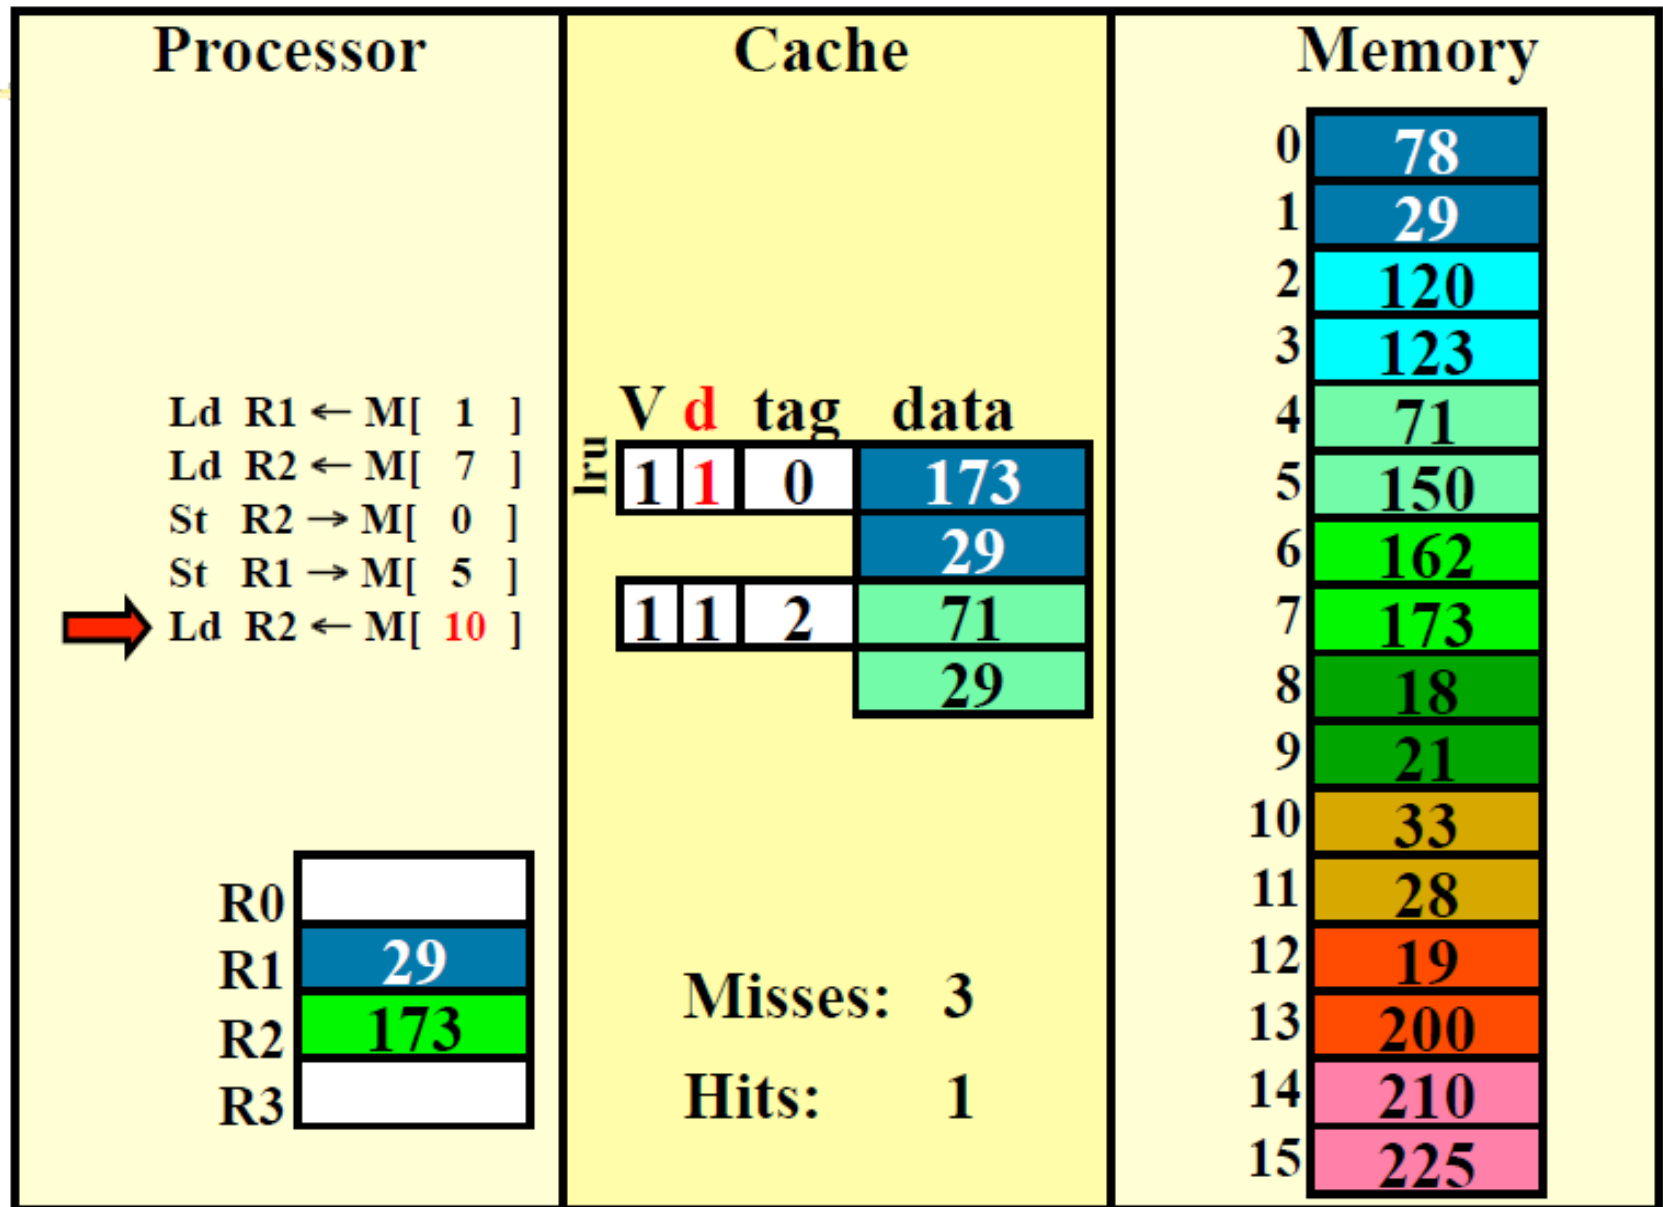
write-back (REF 4)



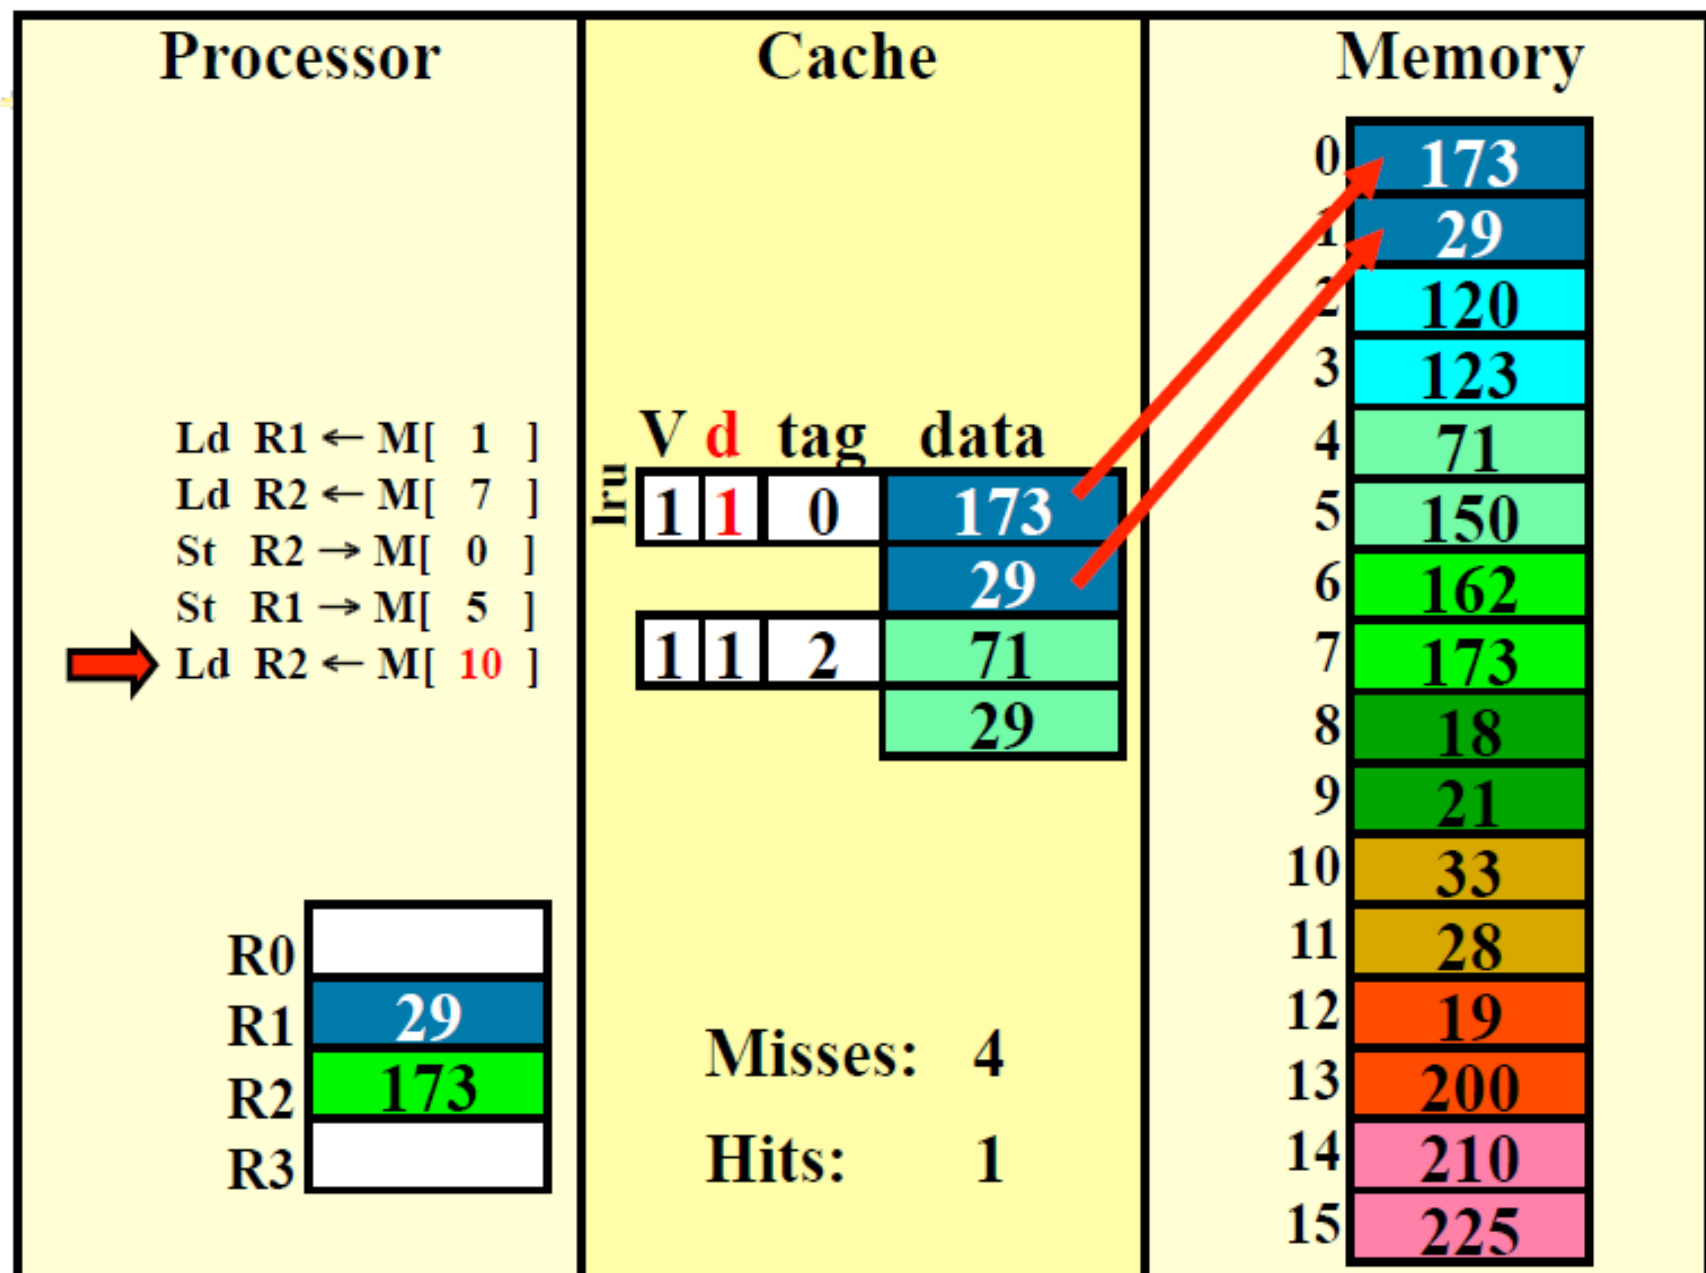
write-back (REF 4)



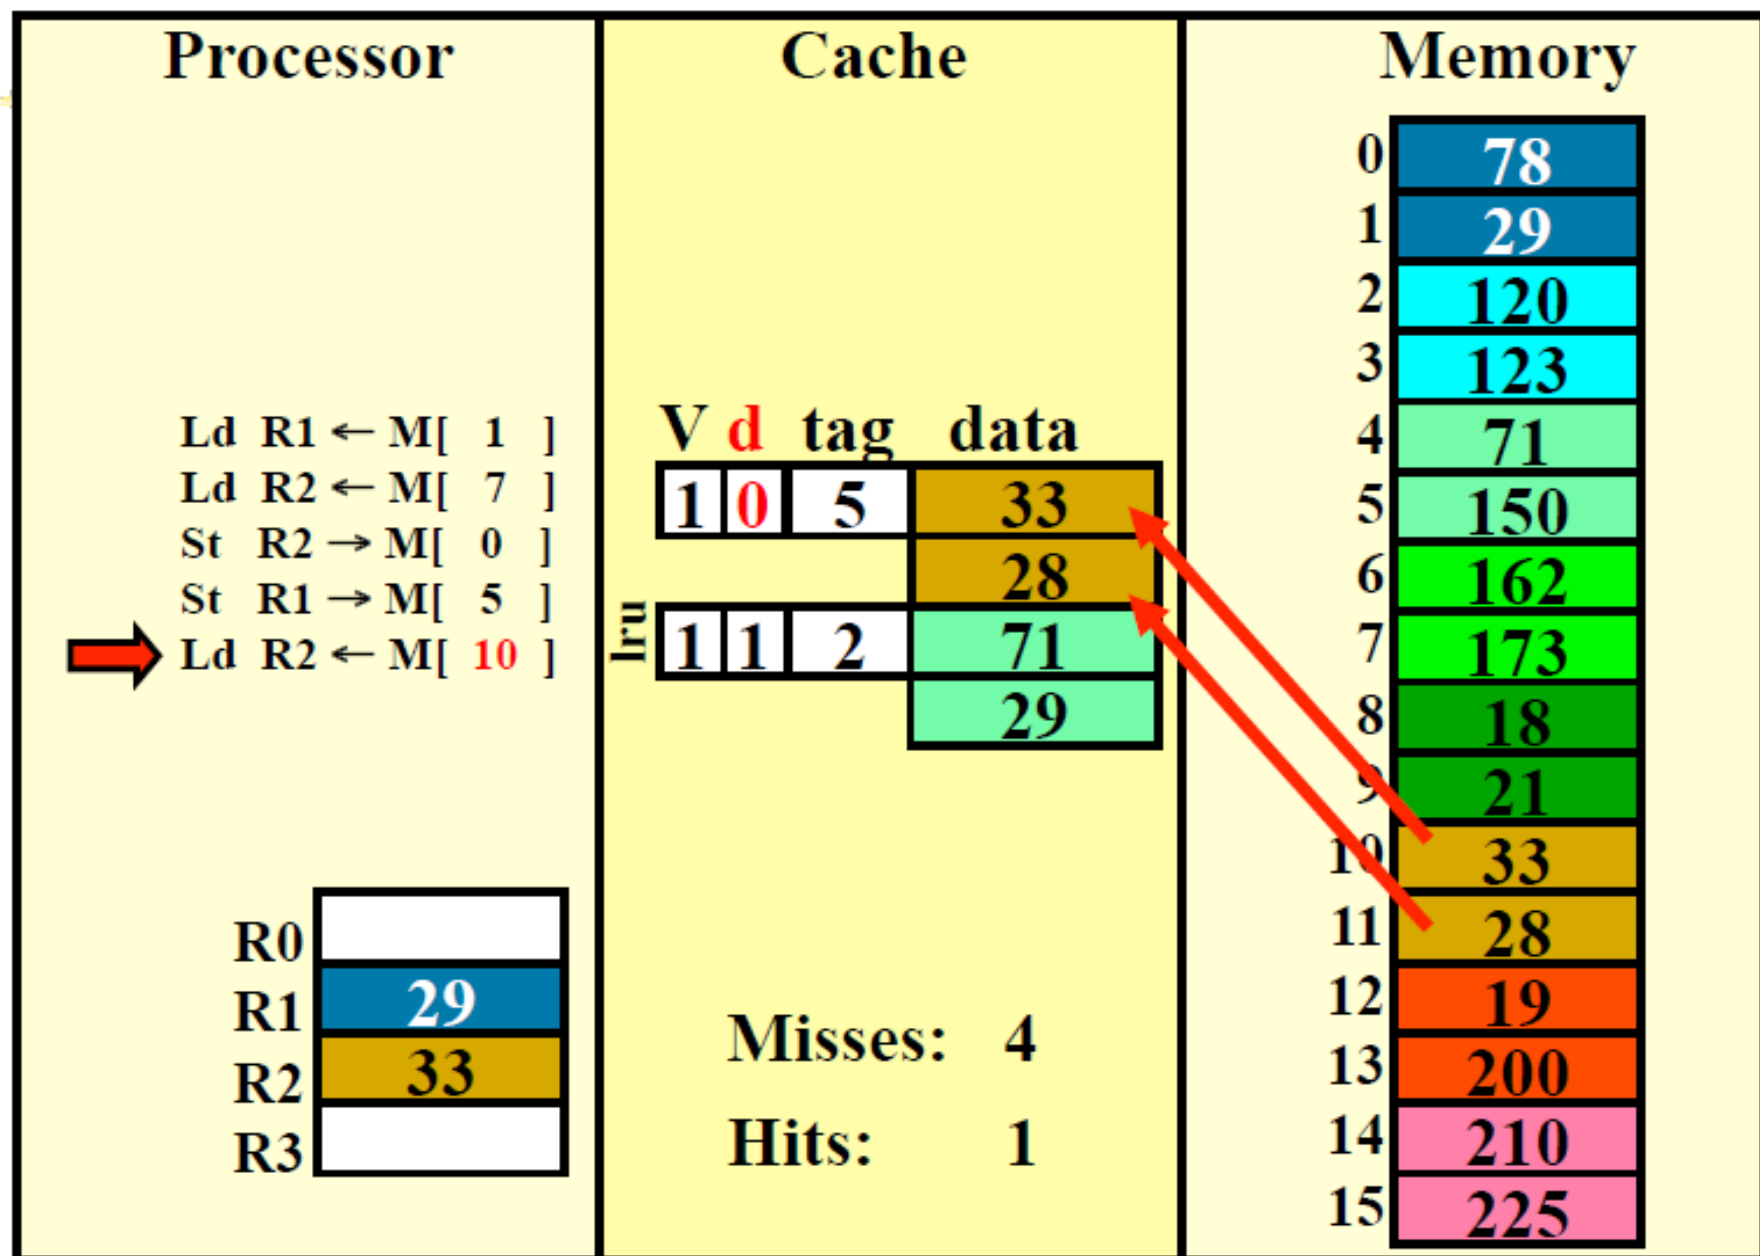
write-back (REF 5)



write-back (REF 5)



write-back (REF 5)



How many memory references?


- Each miss reads a block
 - 2 bytes in this cache
- Each evicted dirty cache line writes a block
- Total reads: 8 bytes
- Total writes: 4 bytes (after final eviction)

Choose write-back or write-through?

VIRTUAL MEMORY

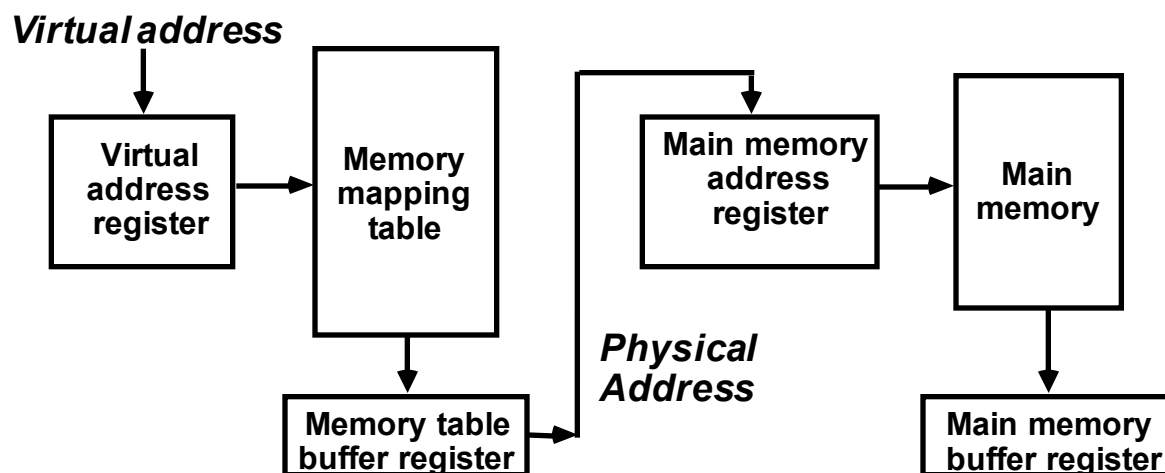
Give the programmer the illusion that the system has a very large memory, even though the computer actually has a relatively small main memory

Address Space(Logical) and Memory Space(Physical)



Address Mapping

Memory Mapping Table for Virtual Address -> Physical Address



ADDRESS MAPPING

Address Space and Memory Space are each divided into fixed size group of words called *blocks* or *pages*

1K words group

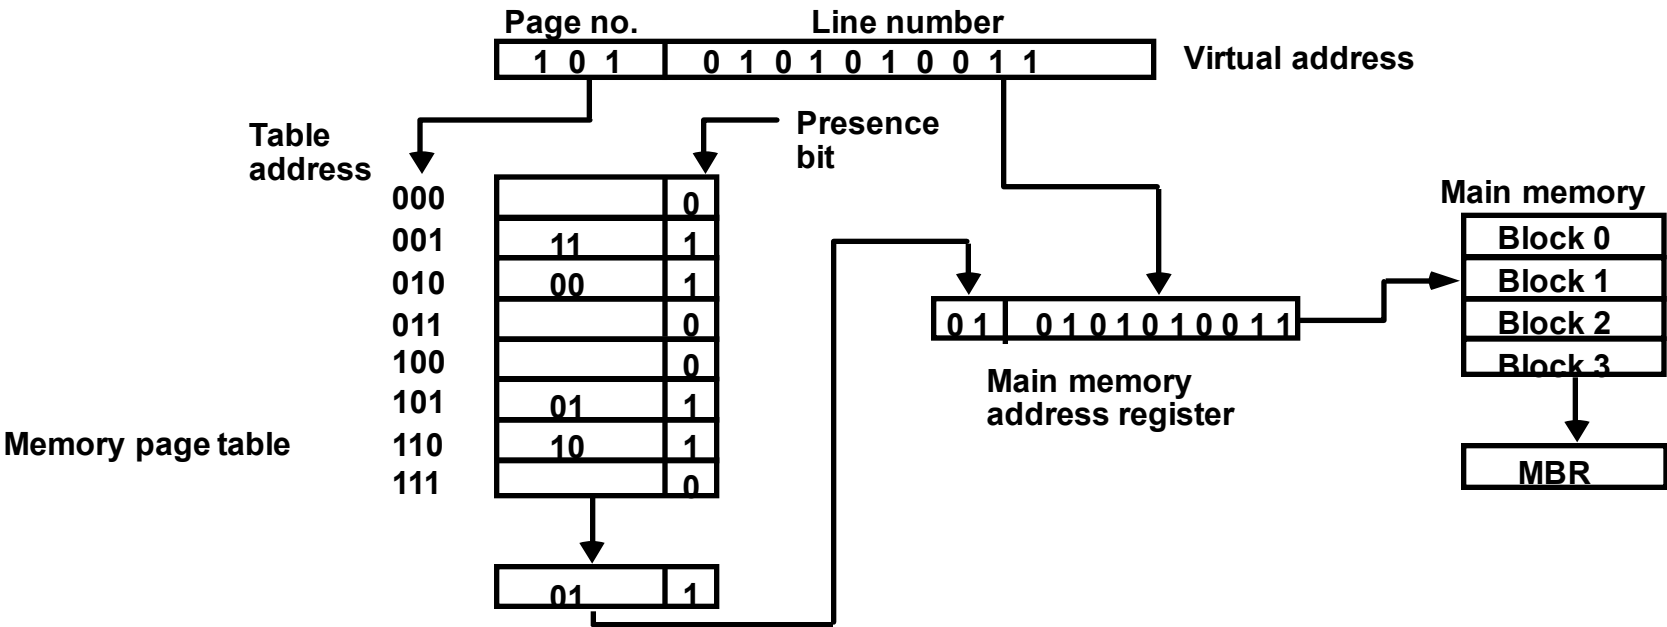
Address space
 $N = 8K = 2^{13}$

Page 0
Page 1
Page 2
Page 3
Page 4
Page 5
Page 6
Page 7

Memory space
 $M = 4K = 2^{12}$

Block 0
Block 1
Block 2
Block 3

Organization of memory Mapping Table in a paged system



ASSOCIATIVE MEMORY PAGE TABLE

Assume that

Number of Blocks in memory = m

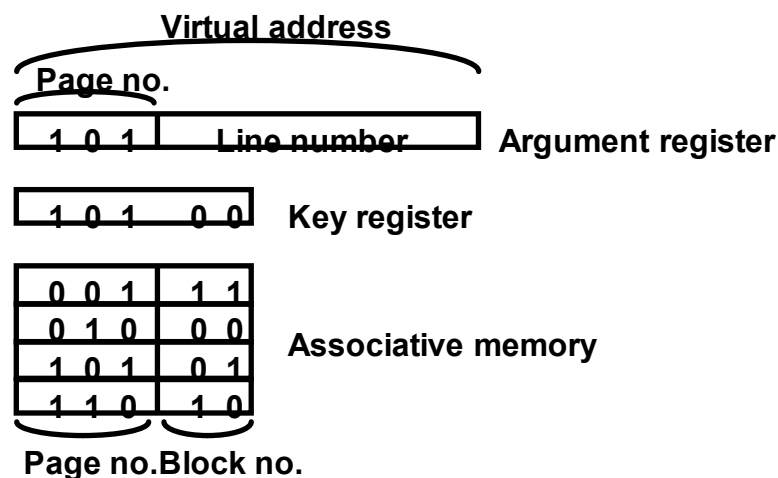
Number of Pages in Virtual Address Space = n

Page Table

- Straight forward design \rightarrow n entry table in memory
Inefficient storage space utilization
 \leftarrow $n-m$ entries of the table is empty

- More efficient method is m -entry Page Table

Page Table made of an Associative Memory
 m words; (Page Number:Block Number)



Page Fault

Page number cannot be found in the Page Table

MEMORY MANAGEMENT HARDWARE


Basic Functions of MM

- *Dynamic Storage Relocation* - mapping logical memory references to physical memory references
- Provision for *Sharing* common information stored in memory by different users
- *Protection* of information against unauthorized access

Segmentation

- A segment is a set of logically related instructions or data elements associated with a given name
- Variable size

User's view of memory




The user does not think of memory as a linear array of words. Rather the user prefers to view memory as a collection of variable sized segments, with no necessary ordering among segments.

User's view of a program

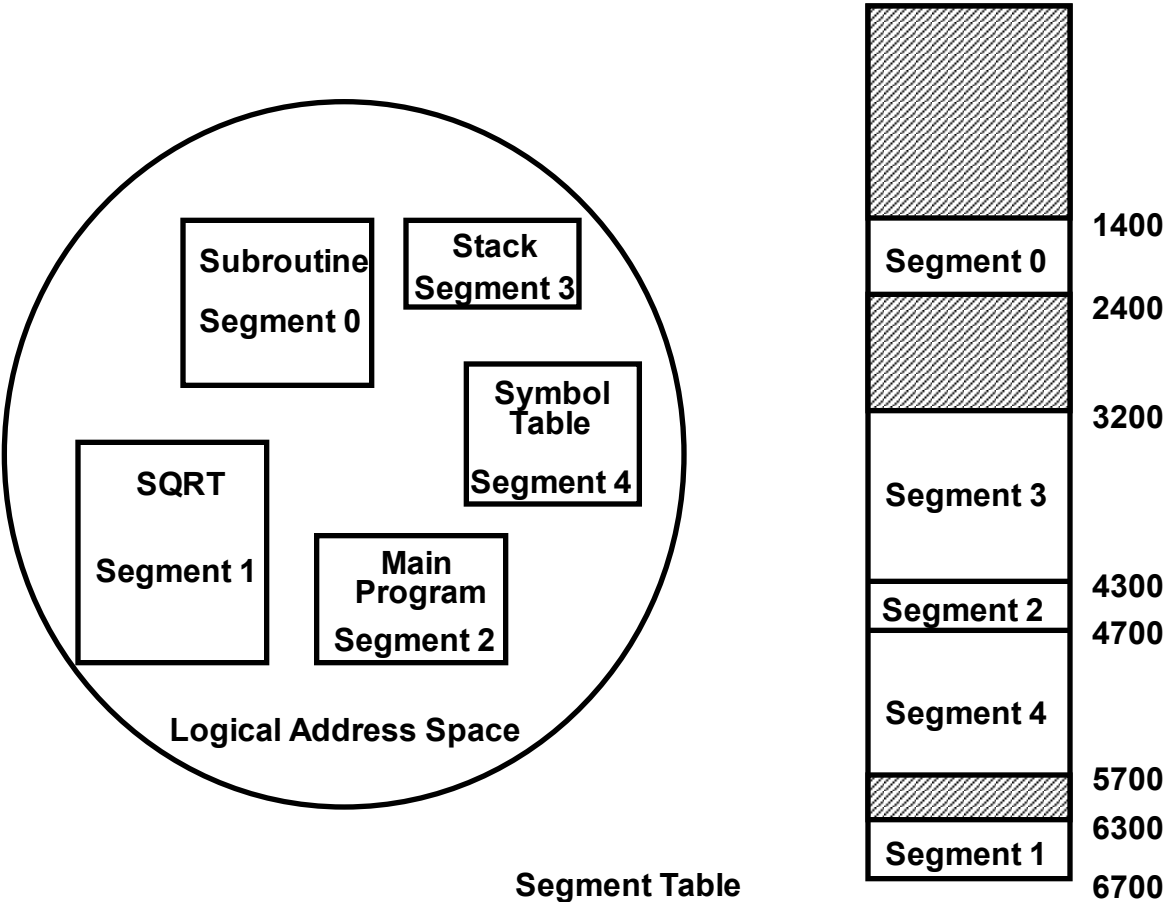
SEGMENTATION

- A memory management scheme which supports user's view of memory
- A logical address space is a collection of segments
- Each segment has a name and a length
- Address specify both the segment name and the offset within the segment.
- For simplicity of implementations, segments are numbered.

Segmentation Hardware

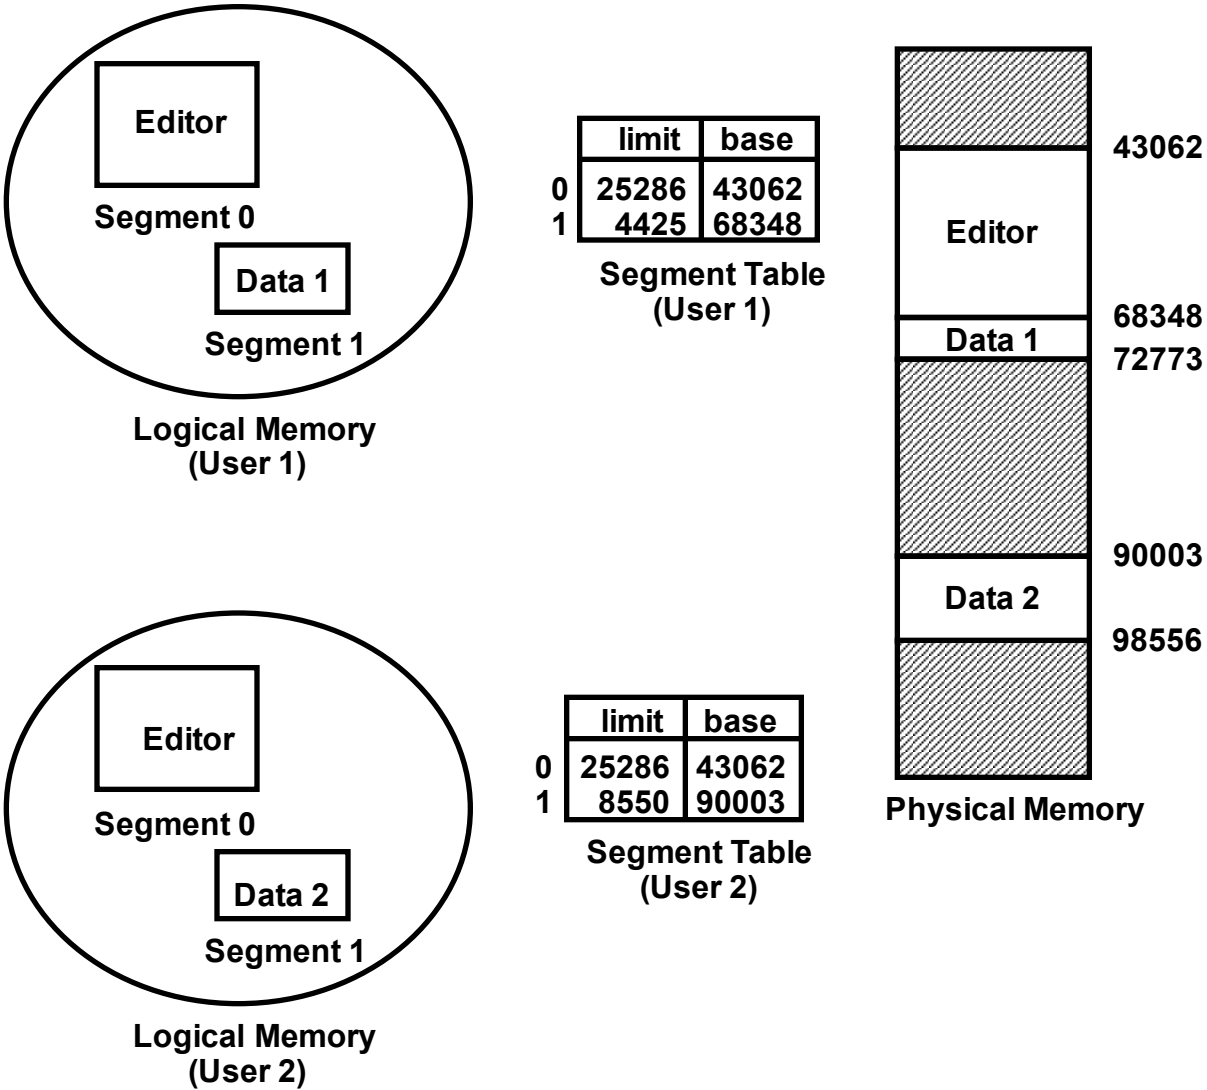


SEGMENTATION EXAMPLE

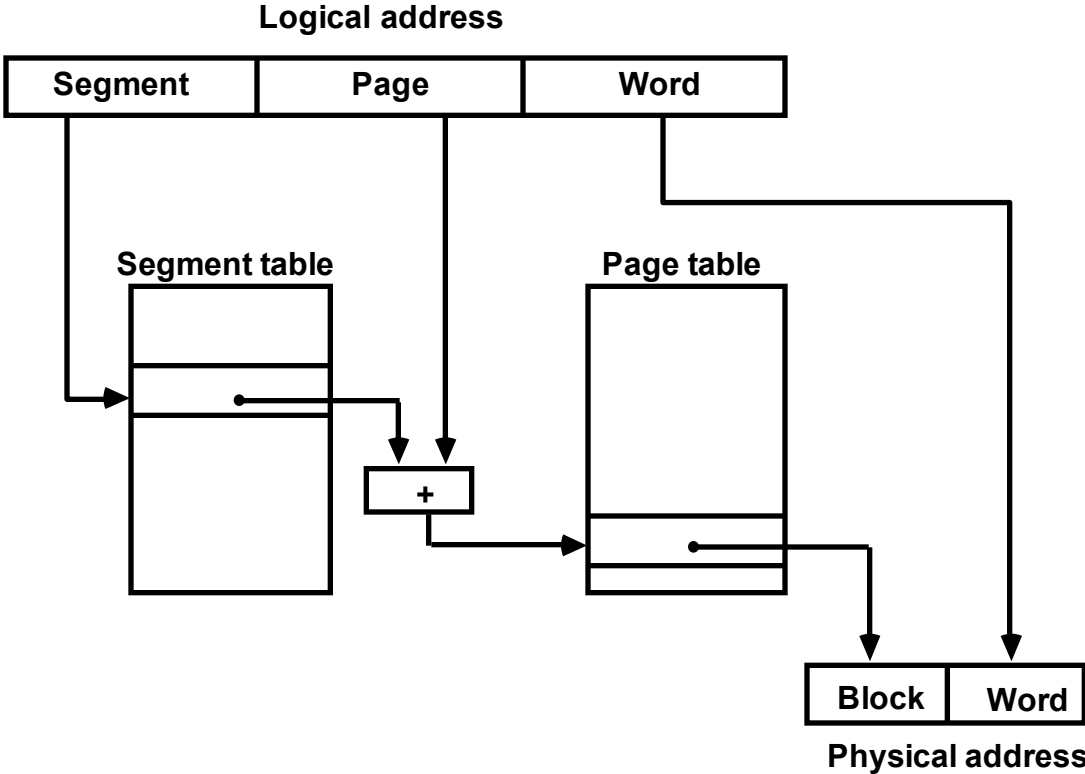


Segment Table		
	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

SHARING OF SEGMENTS



SEGMENTED PAGE SYSTEM



IMPLEMENTATION OF PAGE AND SEGMENT TABLES

Implementation of the Page Table

- Hardware registers (if the page table is reasonably small)
- Main memory
 - Page Table Base Register(PTBR) points to PT
 - Two memory accesses are needed to access a word; one for the page table, one for the word
- Cache memory
 - To speedup the effective memory access time, a special small memory called associative memory, or cache is used

Implementation of the Segment Table

Similar to the case of the page table

EXAMPLE

Logical and Physical Addresses

Logical address format: 16 segments of 256 pages each, each page has 256words

4	8	8
Segment	Page	Word

2²⁰ x 32
Physical
memory

Physical address format: 4096 blocks of 256 words each, each word has 32bits

12	8
Block	Word

Logical and Physical Memory Address Assignment

Hexa address	Page number
60000	Page 0
60100	Page 1
60200	Page 2
60300	Page 3
60400	Page 4
604FF	

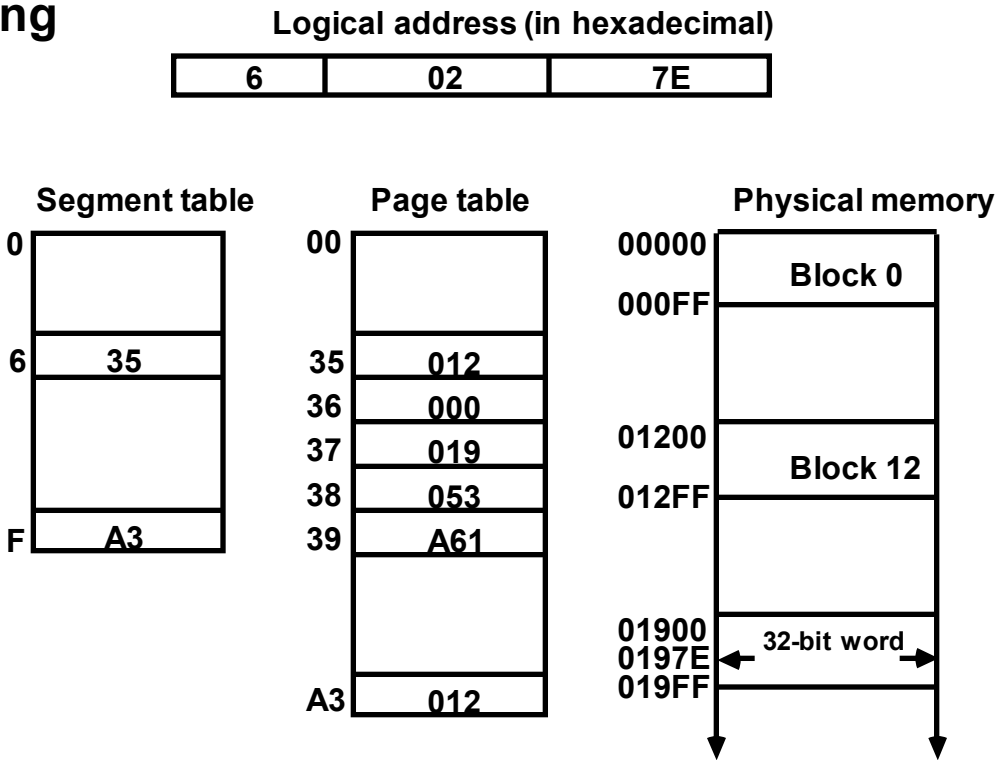
(a) Logical address assignment

Segment	Page	Block
6	00	012
6	01	000
6	02	019
6	03	053
6	04	A61

(b) Segment-page versus
memory block assignment

LOGICAL TO PHYSICAL MEMORY MAPPING

Segment and page table mapping



Associative memory mapping

Segment	Page	Block
6	02	019
6	04	A61